

The Three Elements of PyTorch

Sebastian Raschka

What is PyTorch?

Tensor library

What is PyTorch?

1

Tensor library

What is PyTorch?

1

Tensor library

2

Automatic
differentiation engine

What is PyTorch?

1

Tensor library

2

Automatic
differentiation engine

What is PyTorch?

3

Deep learning
library

Sounds like TensorFlow?

Sounds like TensorFlow?

Yes, kind of.

Python is also just a scripting language

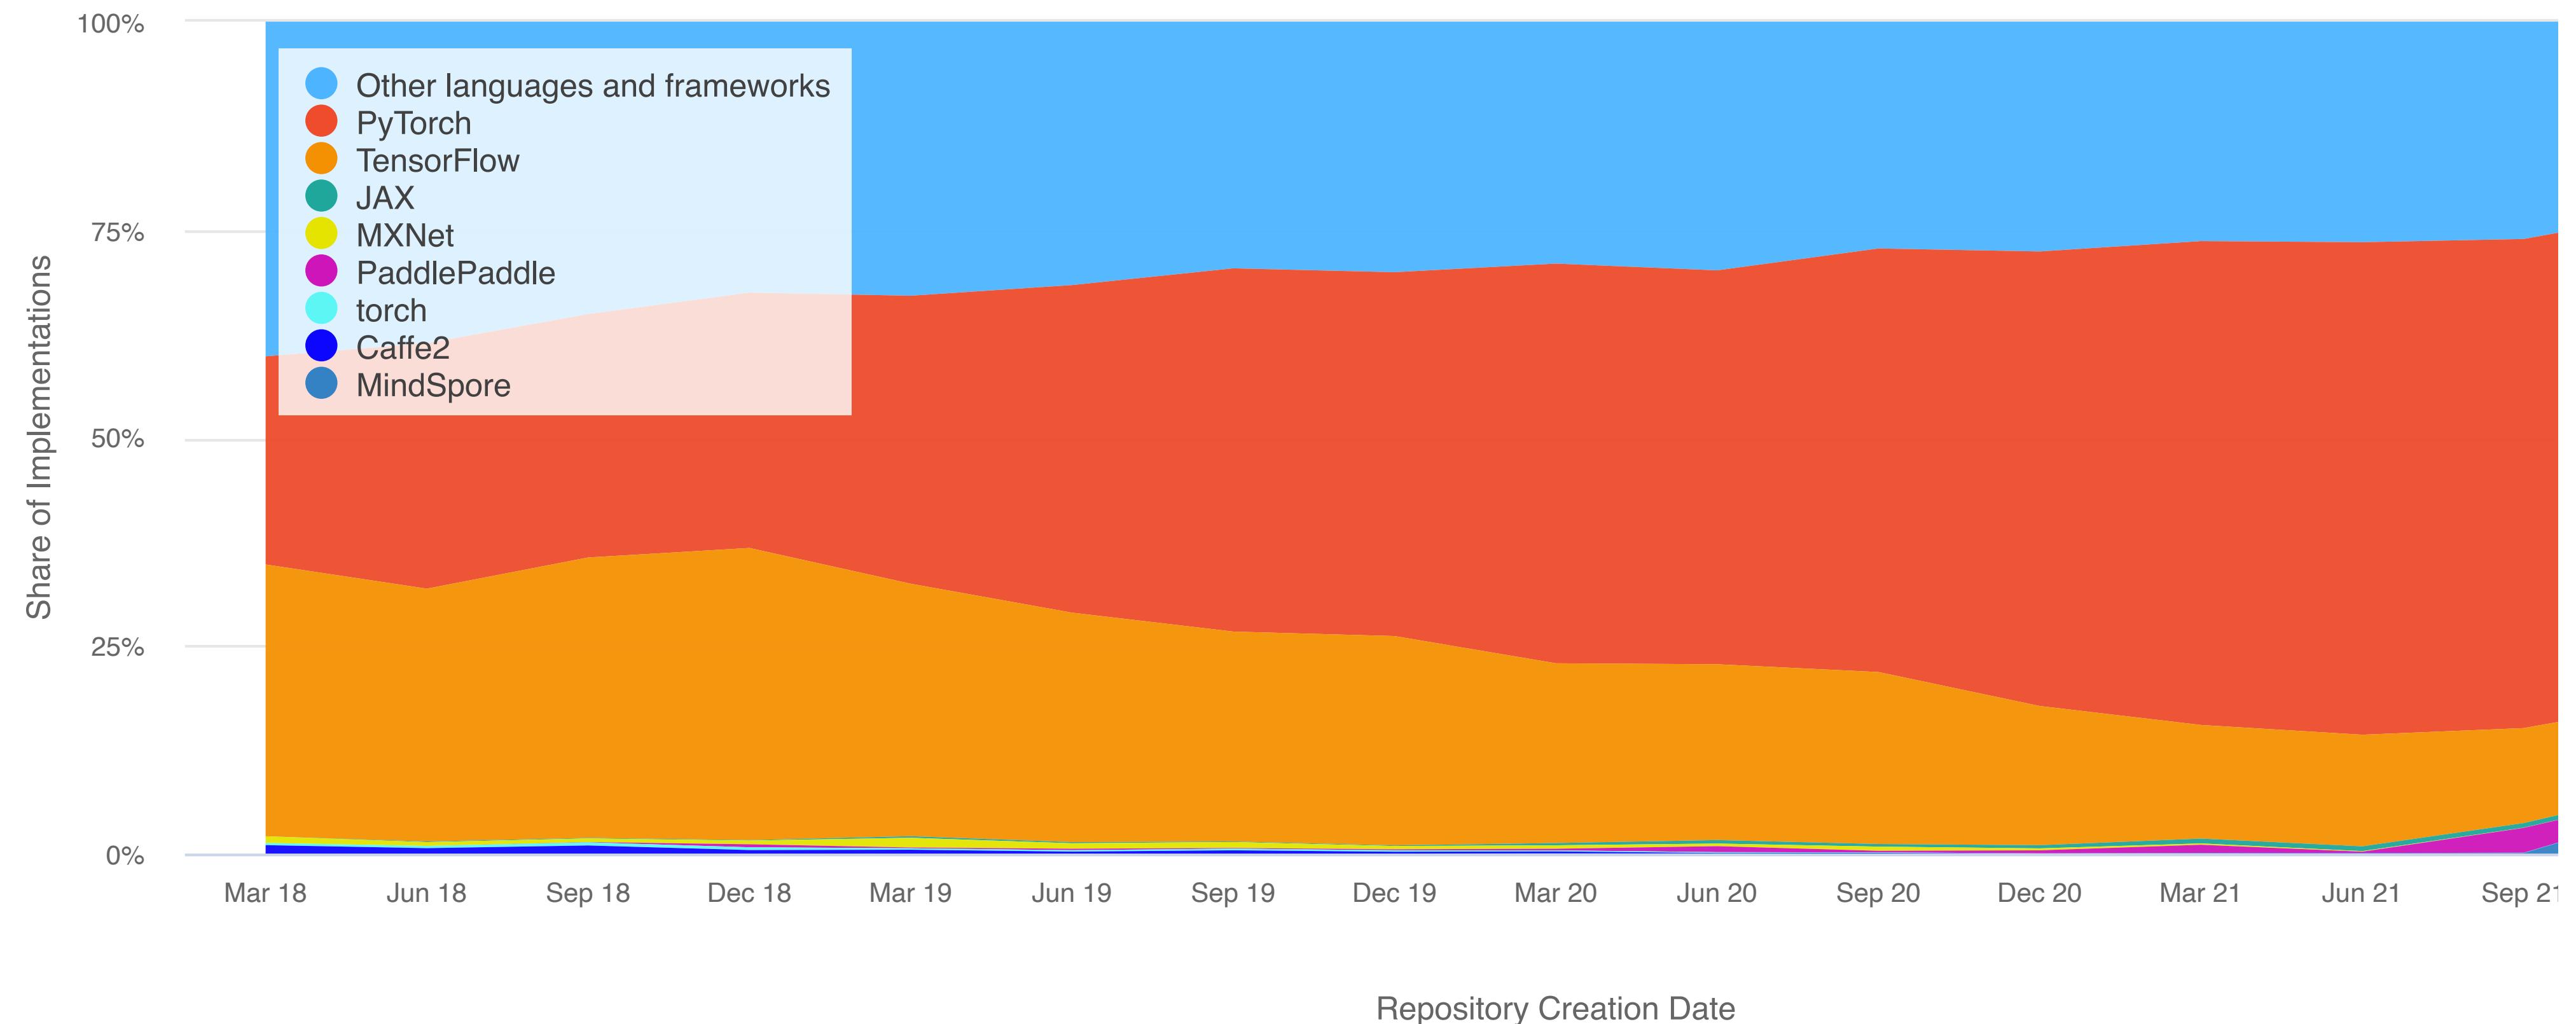
**Python is also just a scripting language
like Ruby or Perl.**

VS Code is also just a text editor

**VS Code is also just a text editor
like VIM or Emacs.**

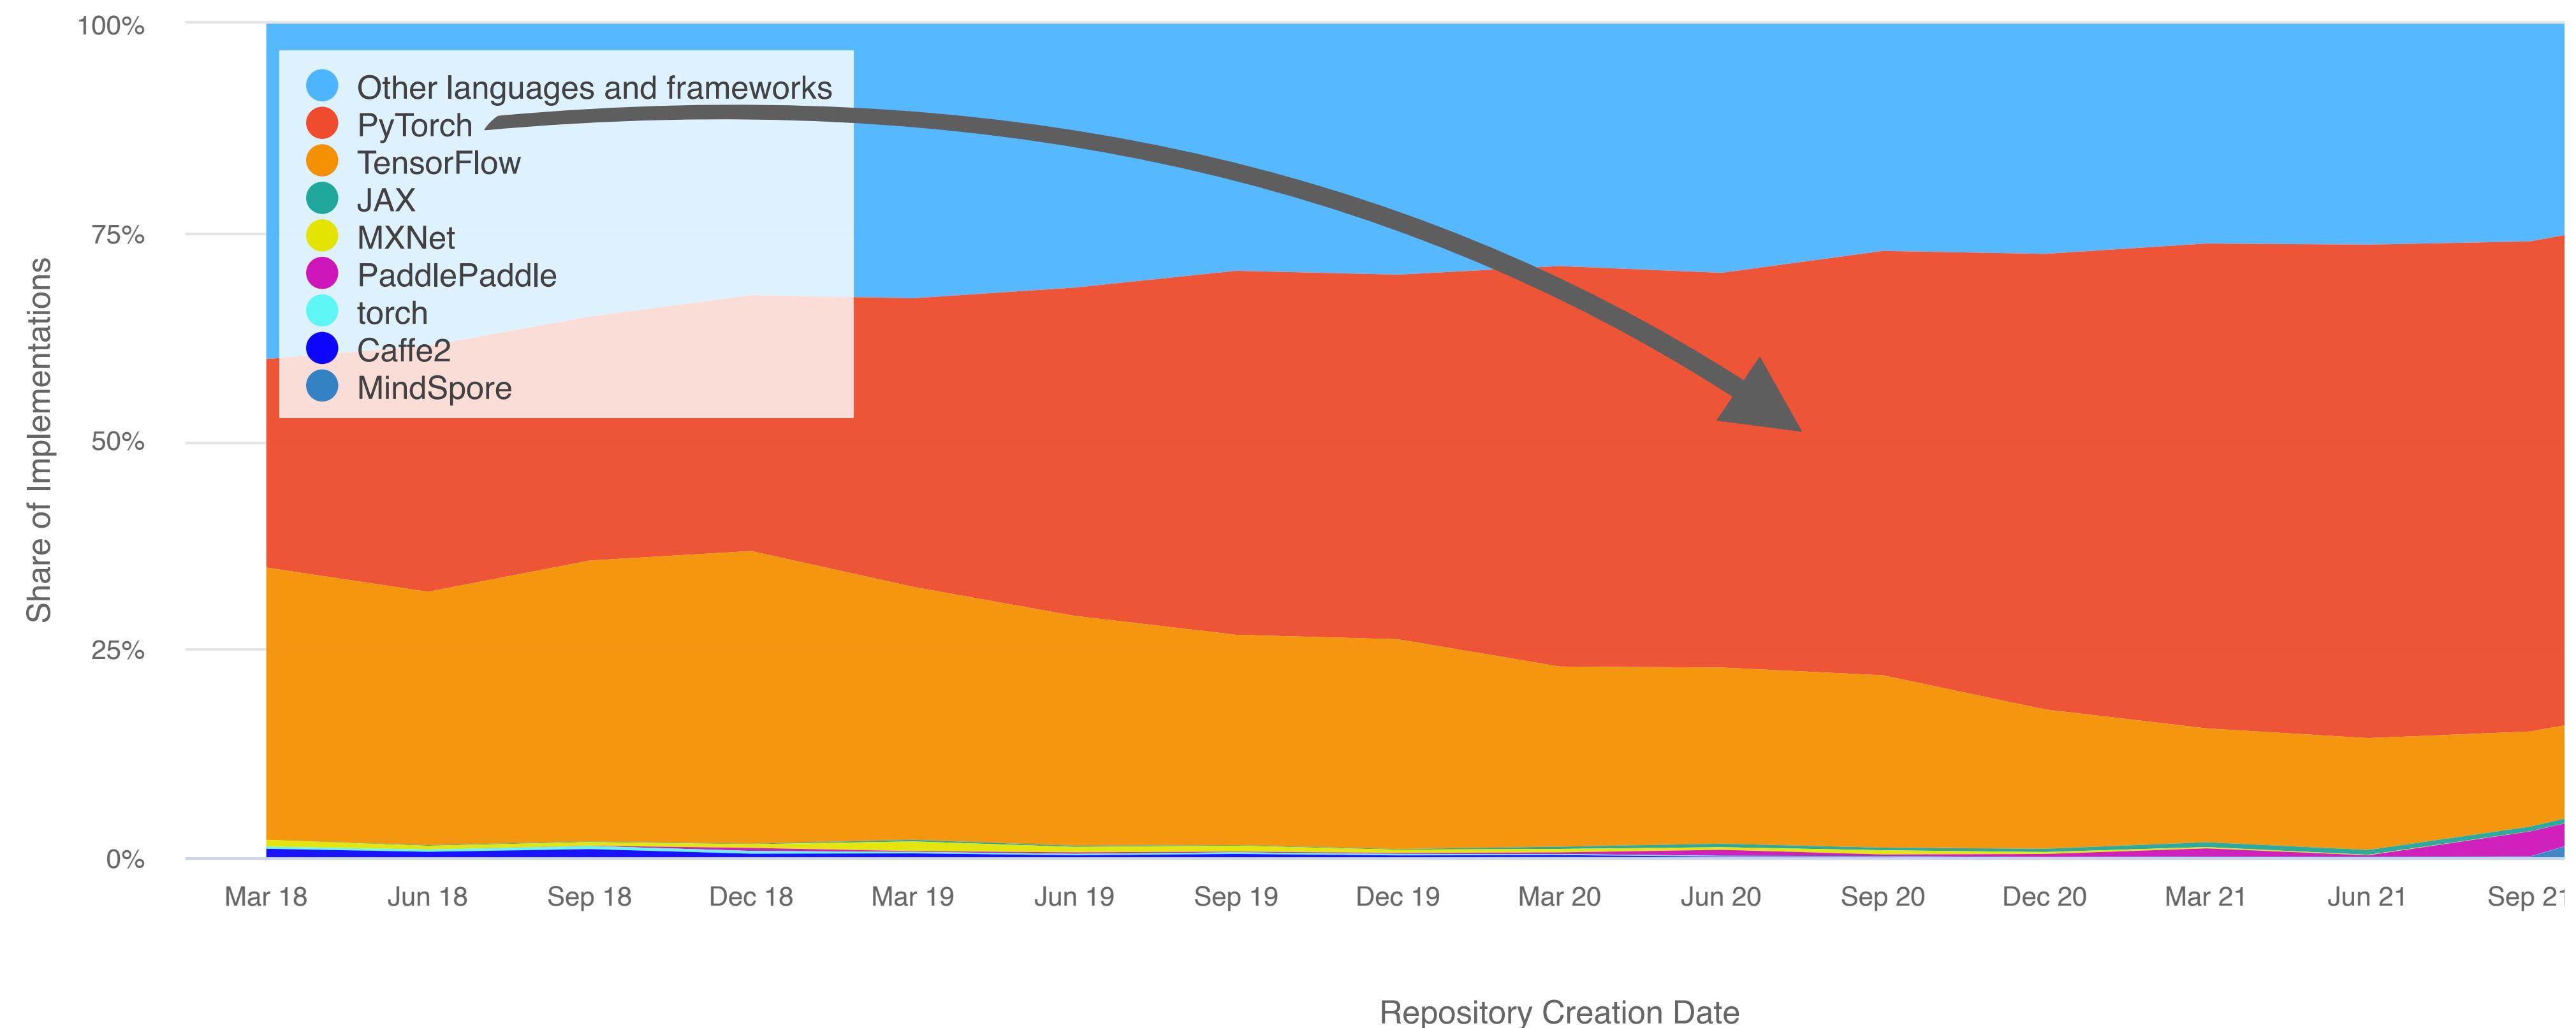
There is probably a reason ...

Paper Implementations grouped by framework



There is probably a reason ...

Paper Implementations grouped by framework



As a former TF user, why do I like PyTorch?

As a former TF user, why do I like PyTorch?

Stay tuned!

1

Tensor library

What is PyTorch?

Tensor library

Scalar (rank-0 tensor)

```
import torch  
  
a = torch.tensor(1.)  
a.shape  
  
torch.Size([])
```

Tensor library

Scalar
(rank-0 tensor)

```
import torch  
  
a = torch.tensor(1.)  
a.shape  
  
torch.Size([])
```

Vector
(rank-1 tensor)

```
a = torch.tensor([1., 2., 3.])  
a.shape  
  
torch.Size([3])
```

Tensor library

Scalar
(rank-0 tensor)

```
import torch
a = torch.tensor(1.)
a.shape
torch.Size([])
```

Vector
(rank-1 tensor)

```
a = torch.tensor([1., 2., 3.])
a.shape
torch.Size([3])
```

Matrix
(rank-2 tensor)

```
a = torch.tensor([[1., 2., 3.],
                  [2., 3., 4.]])
a.shape
torch.Size([2, 3])
```

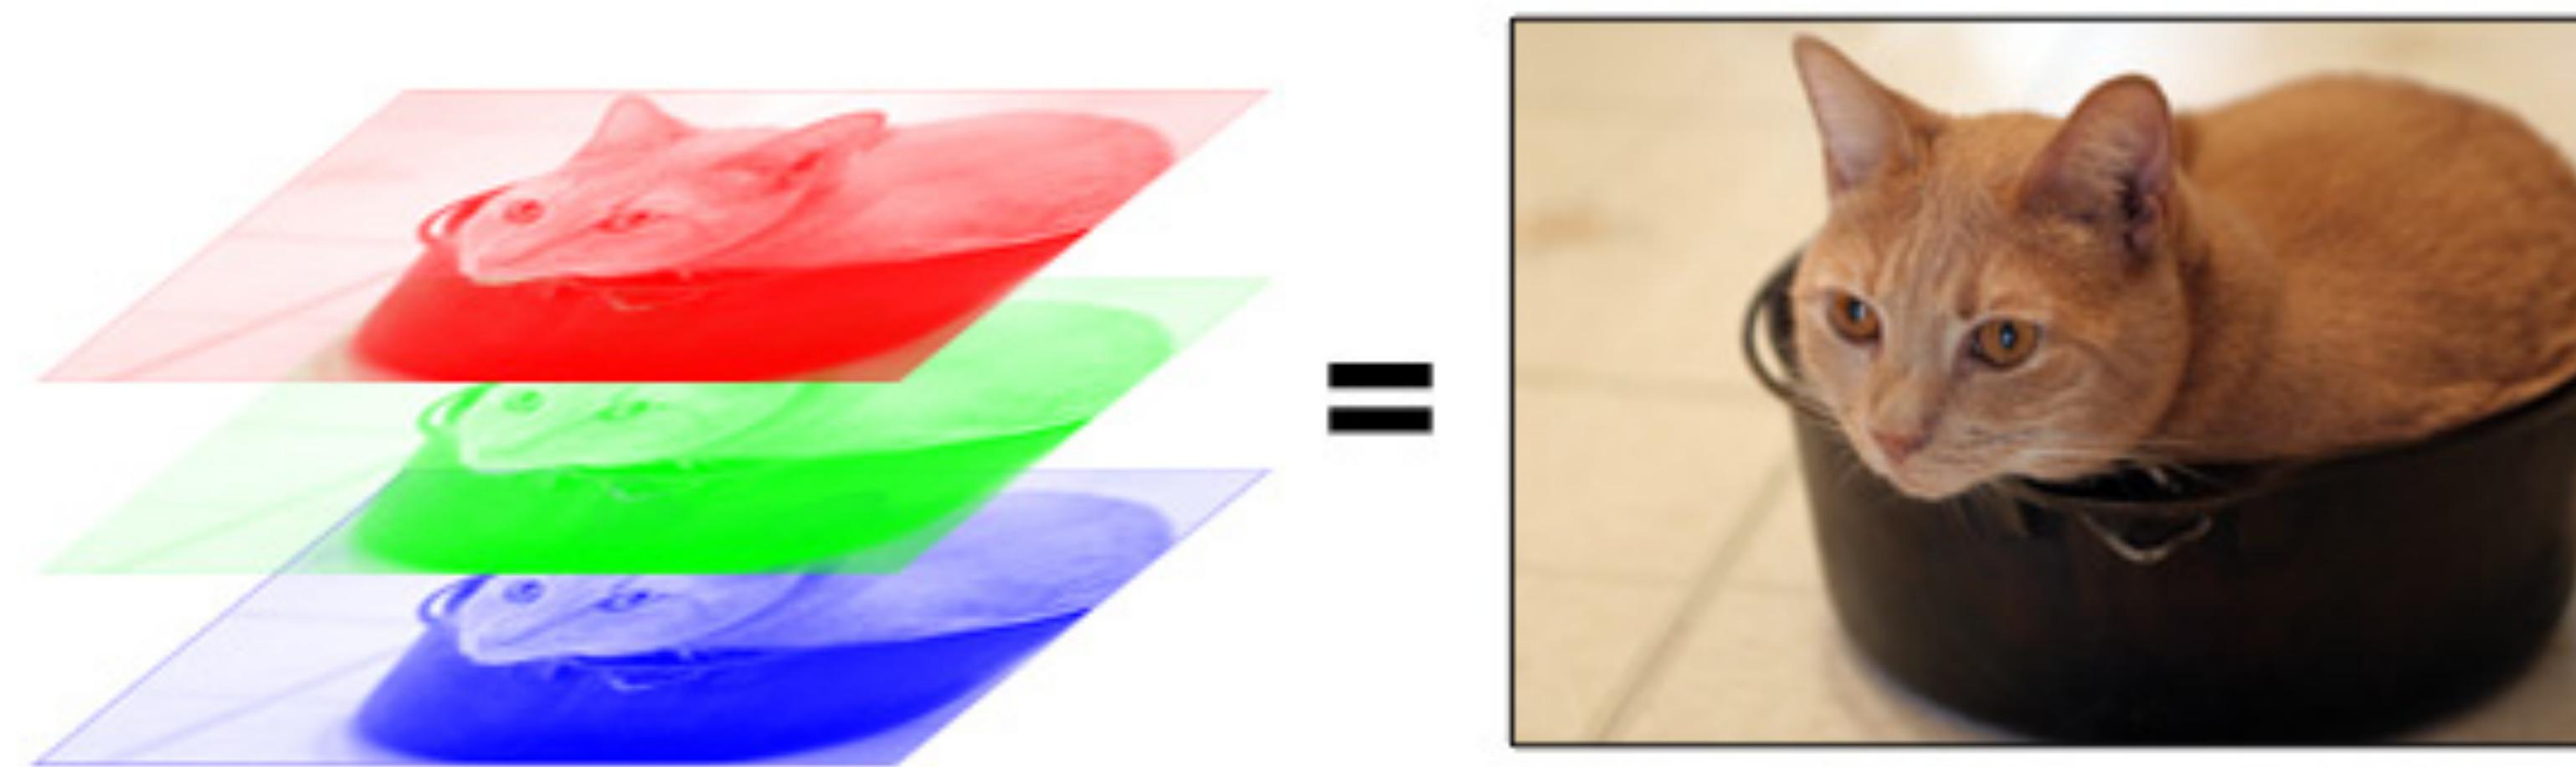
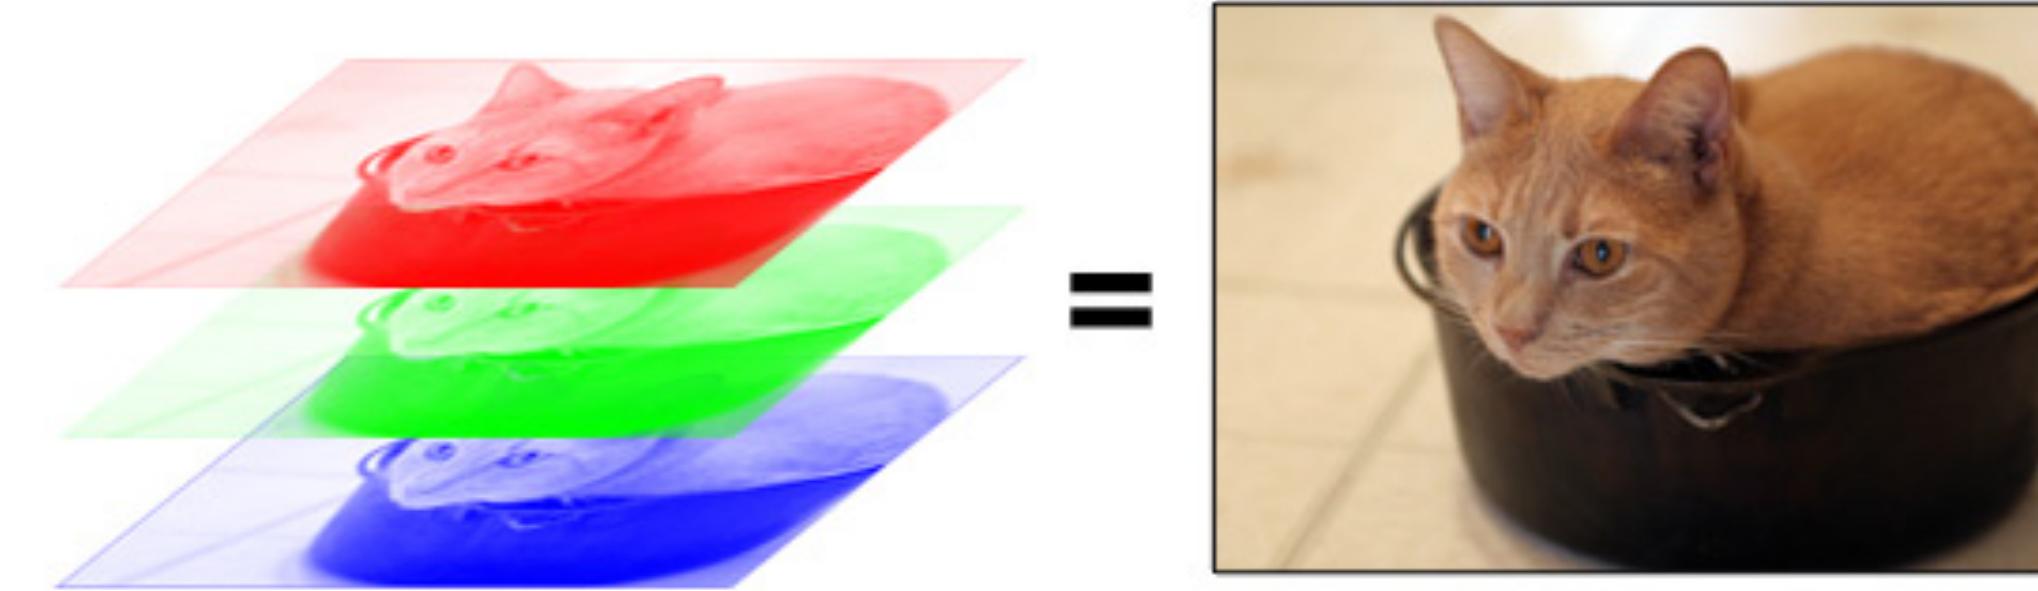


Image Source: <https://code.tutsplus.com/tutorials/create-a-retro-crt-distortion-effect-using-rgb-shifting--active-3359>

Color image as a stack of matrices



Color image as a stack of matrices

3D tensor
(rank-3 tensor)

```
a = torch.tensor([[[1., 2., 3.],  
                  [2., 3., 4.]],  
                  [[5., 6., 7.],  
                   [8., 9., 10.]]])  
a.shape
```

```
torch.Size([2, 2, 3])
```



A stack of multiple color images



A stack of multiple color images

4D tensor
(rank-4 tensor)

```
b = torch.stack((a, a))
b.shape
torch.Size([2, 2, 2, 3])

b
tensor([[[[ 1.,  2.,  3.],
           [ 2.,  3.,  4.]],

          [[ 5.,  6.,  7.],
           [ 8.,  9., 10.]],

          [[[ 1.,  2.,  3.],
            [ 2.,  3.,  4.]],

           [[ 5.,  6.,  7.],
            [ 8.,  9., 10.]]]])
```

Tensor library

Tensor library == array library

Tensor library

`torch.tensor` ≈ `numpy.array`

Tensor library

torch.tensor \approx numpy.array

+ GPU support `a.to('cuda:0')`

Tensor library

`torch.tensor ≈ numpy.array`

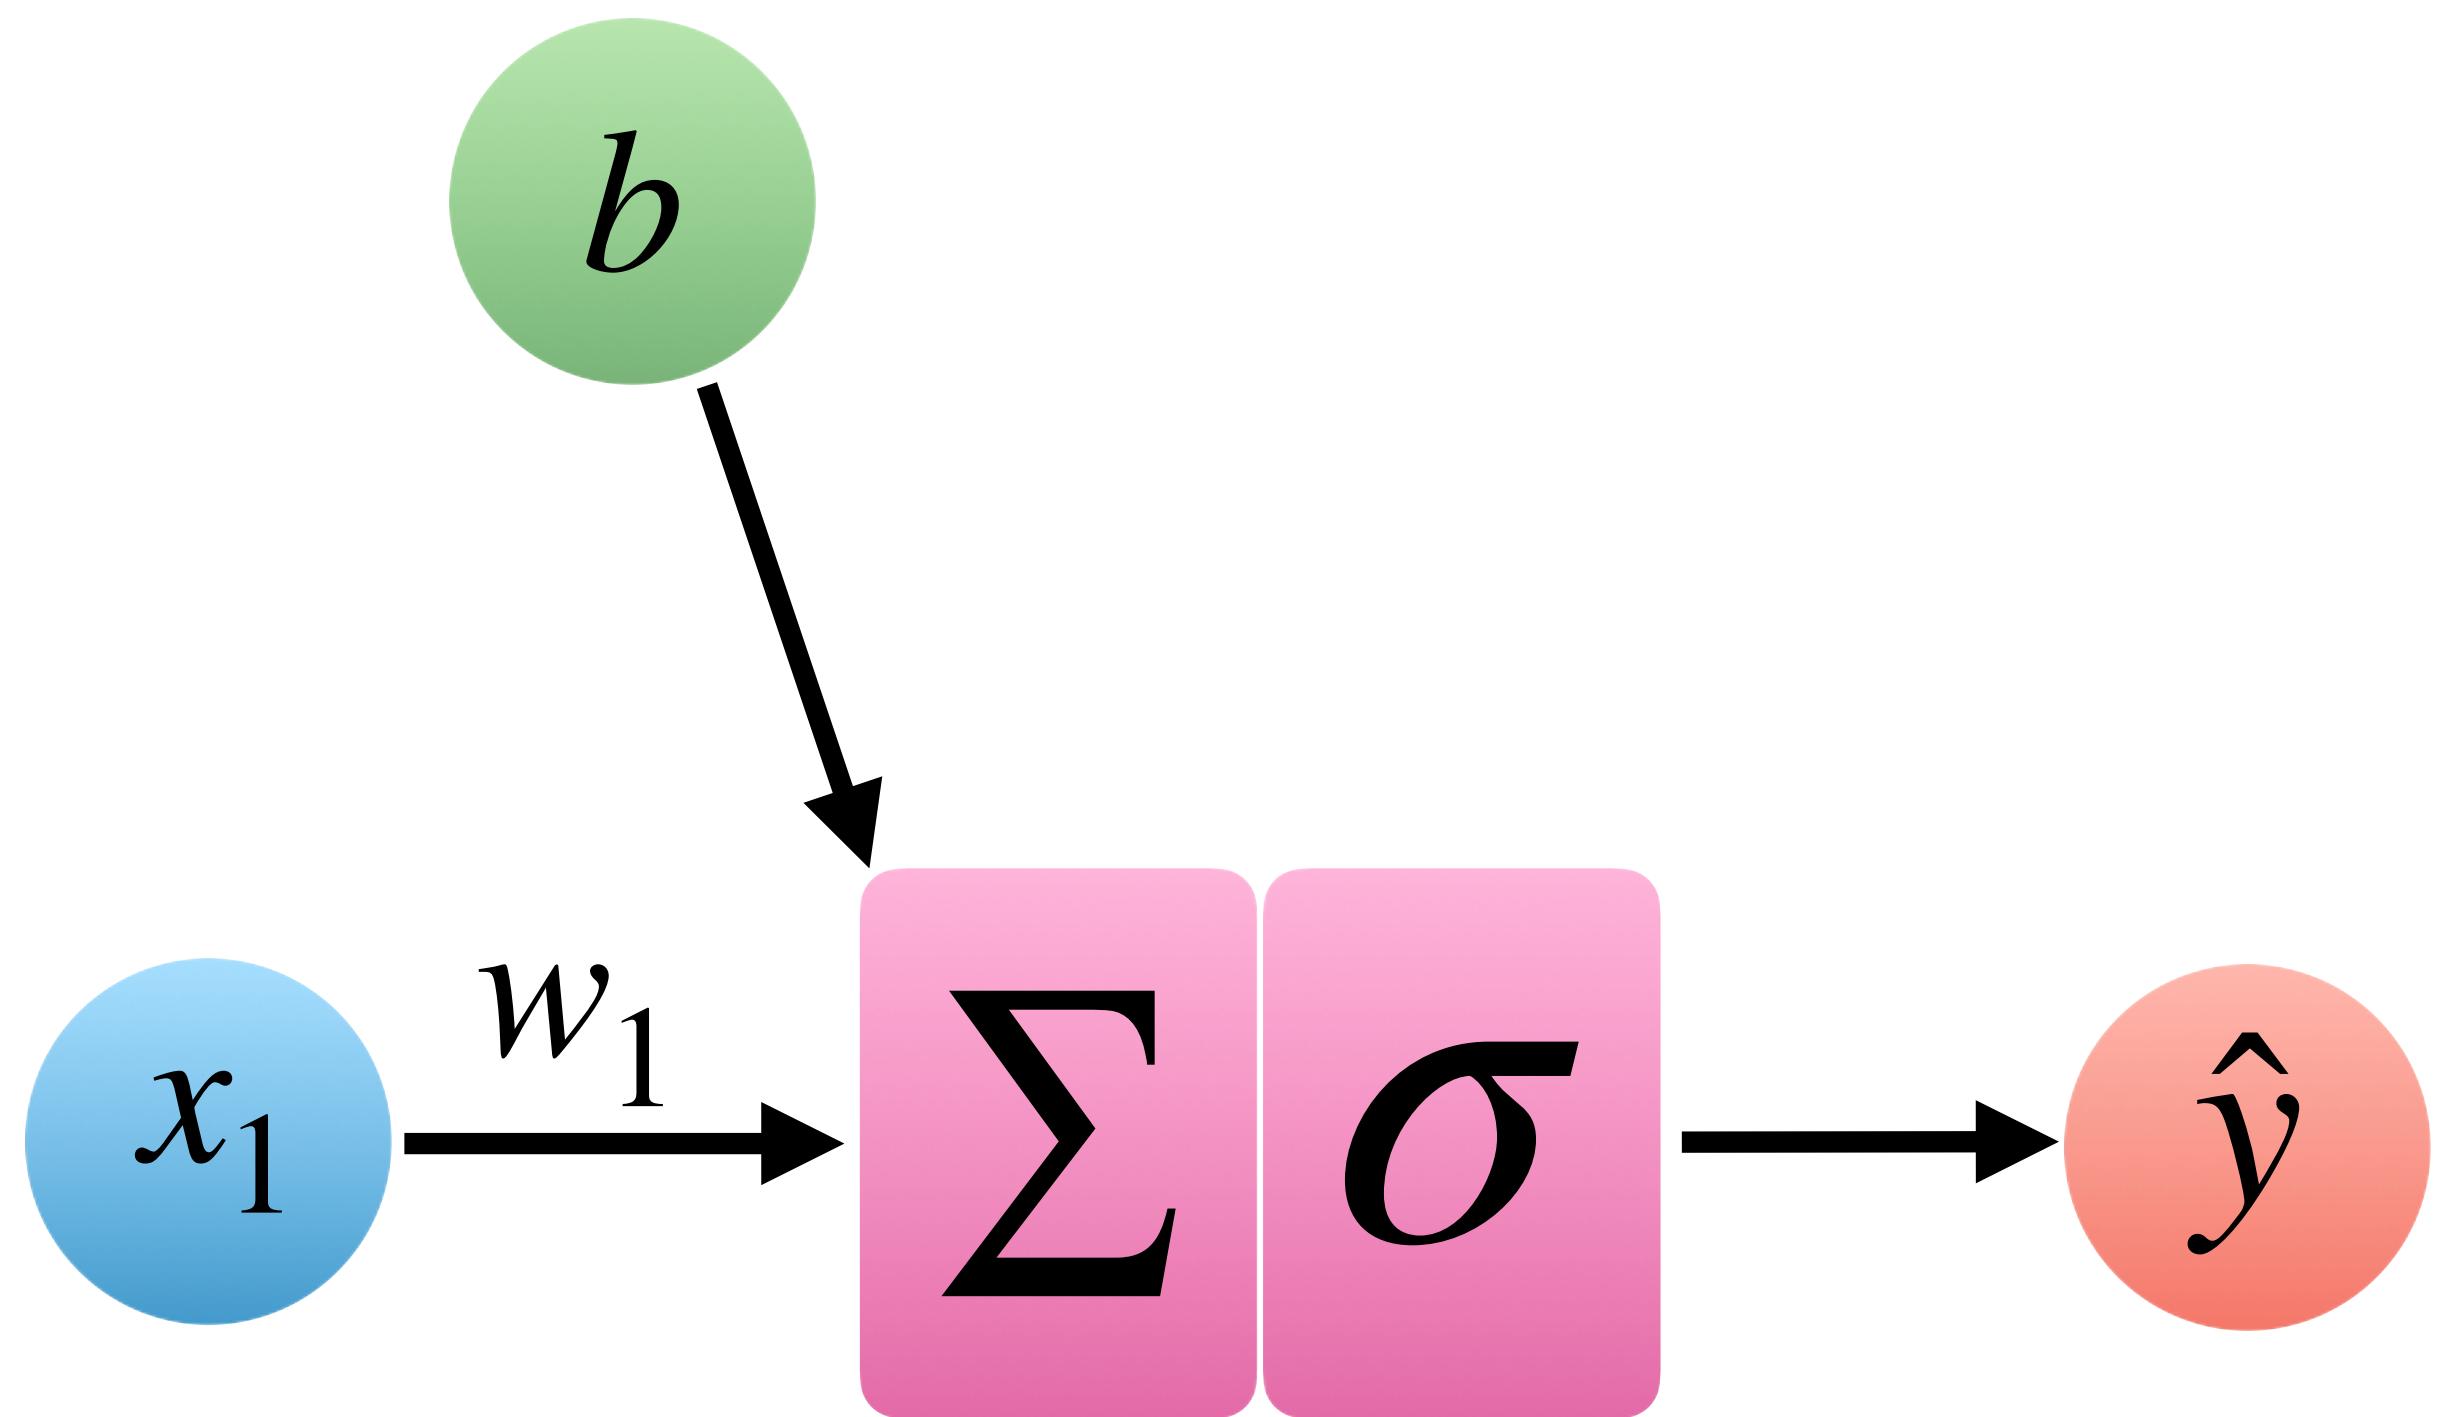
+ GPU support `a.to('cuda:0')`

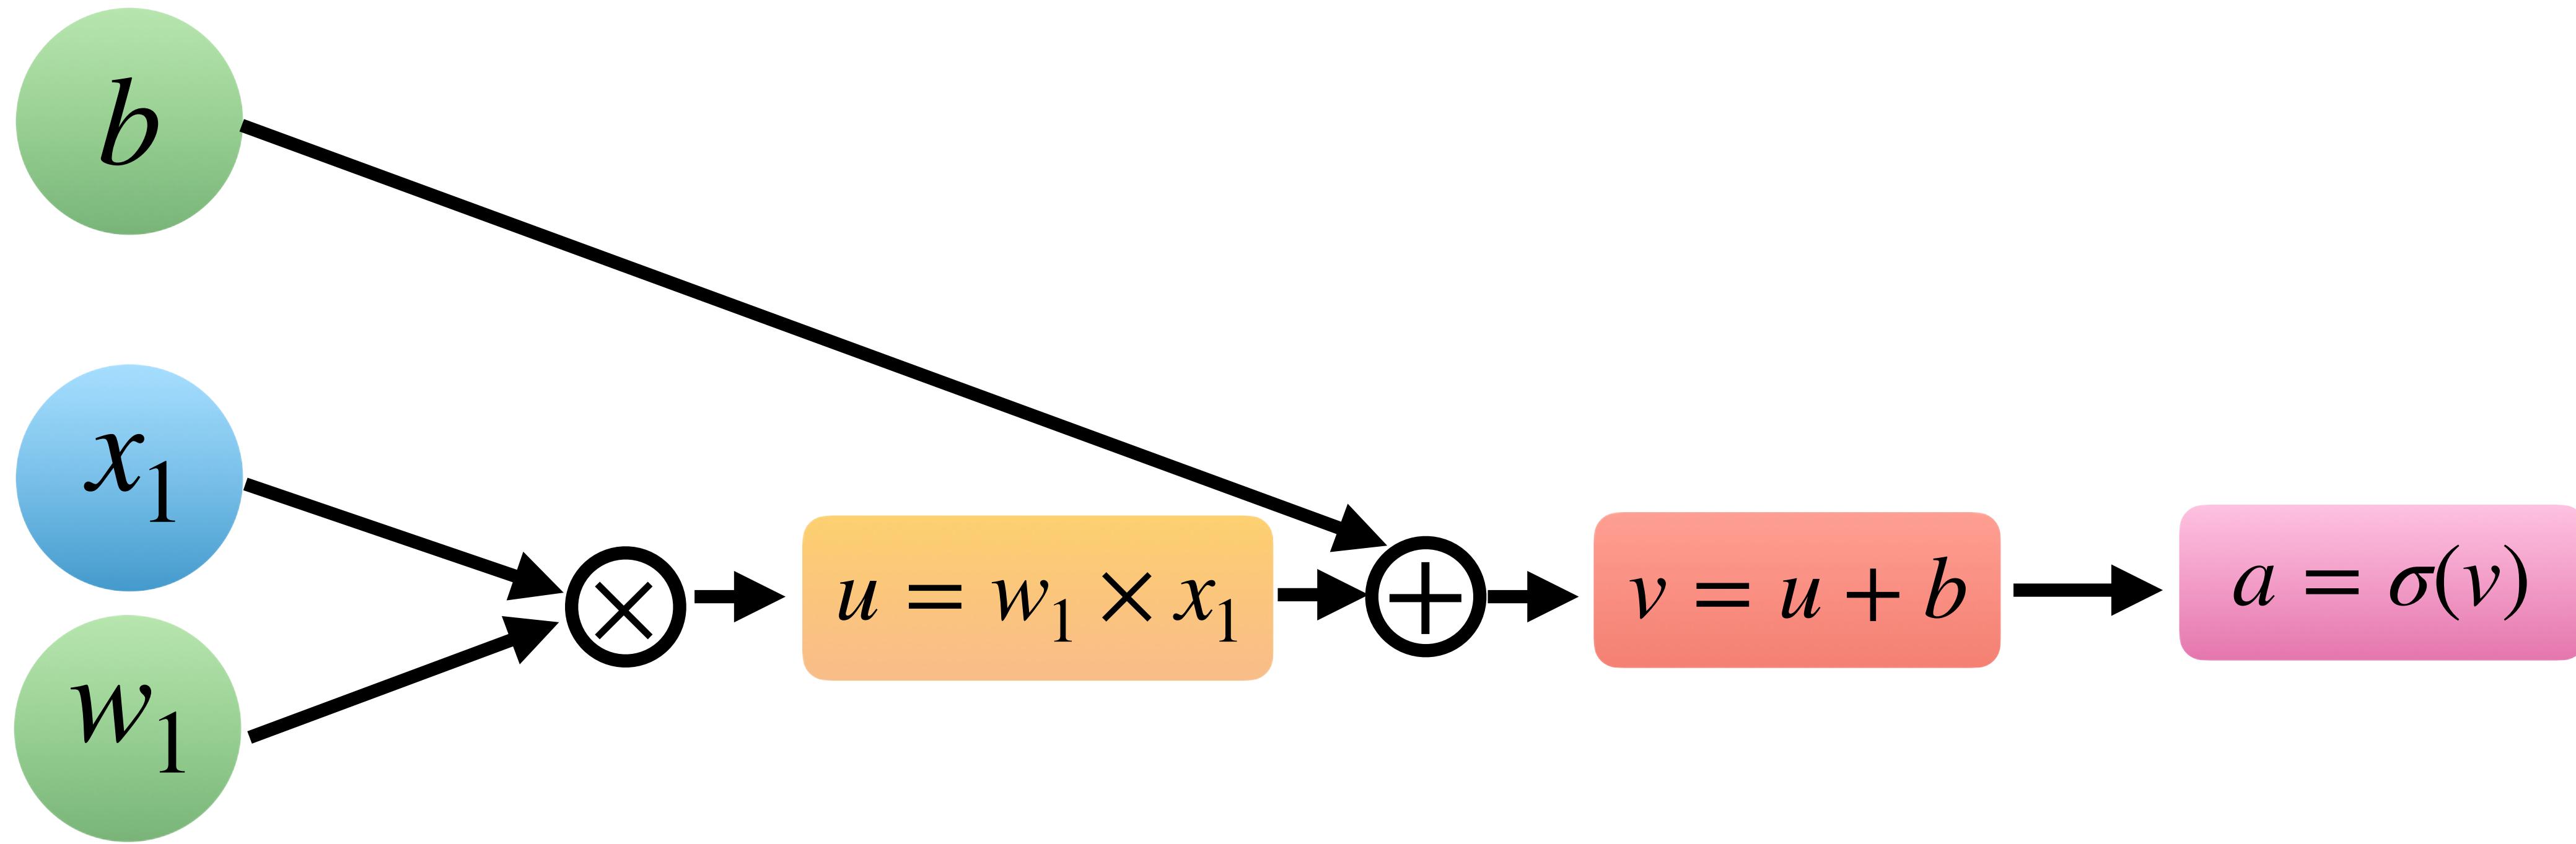
+ autodiff support

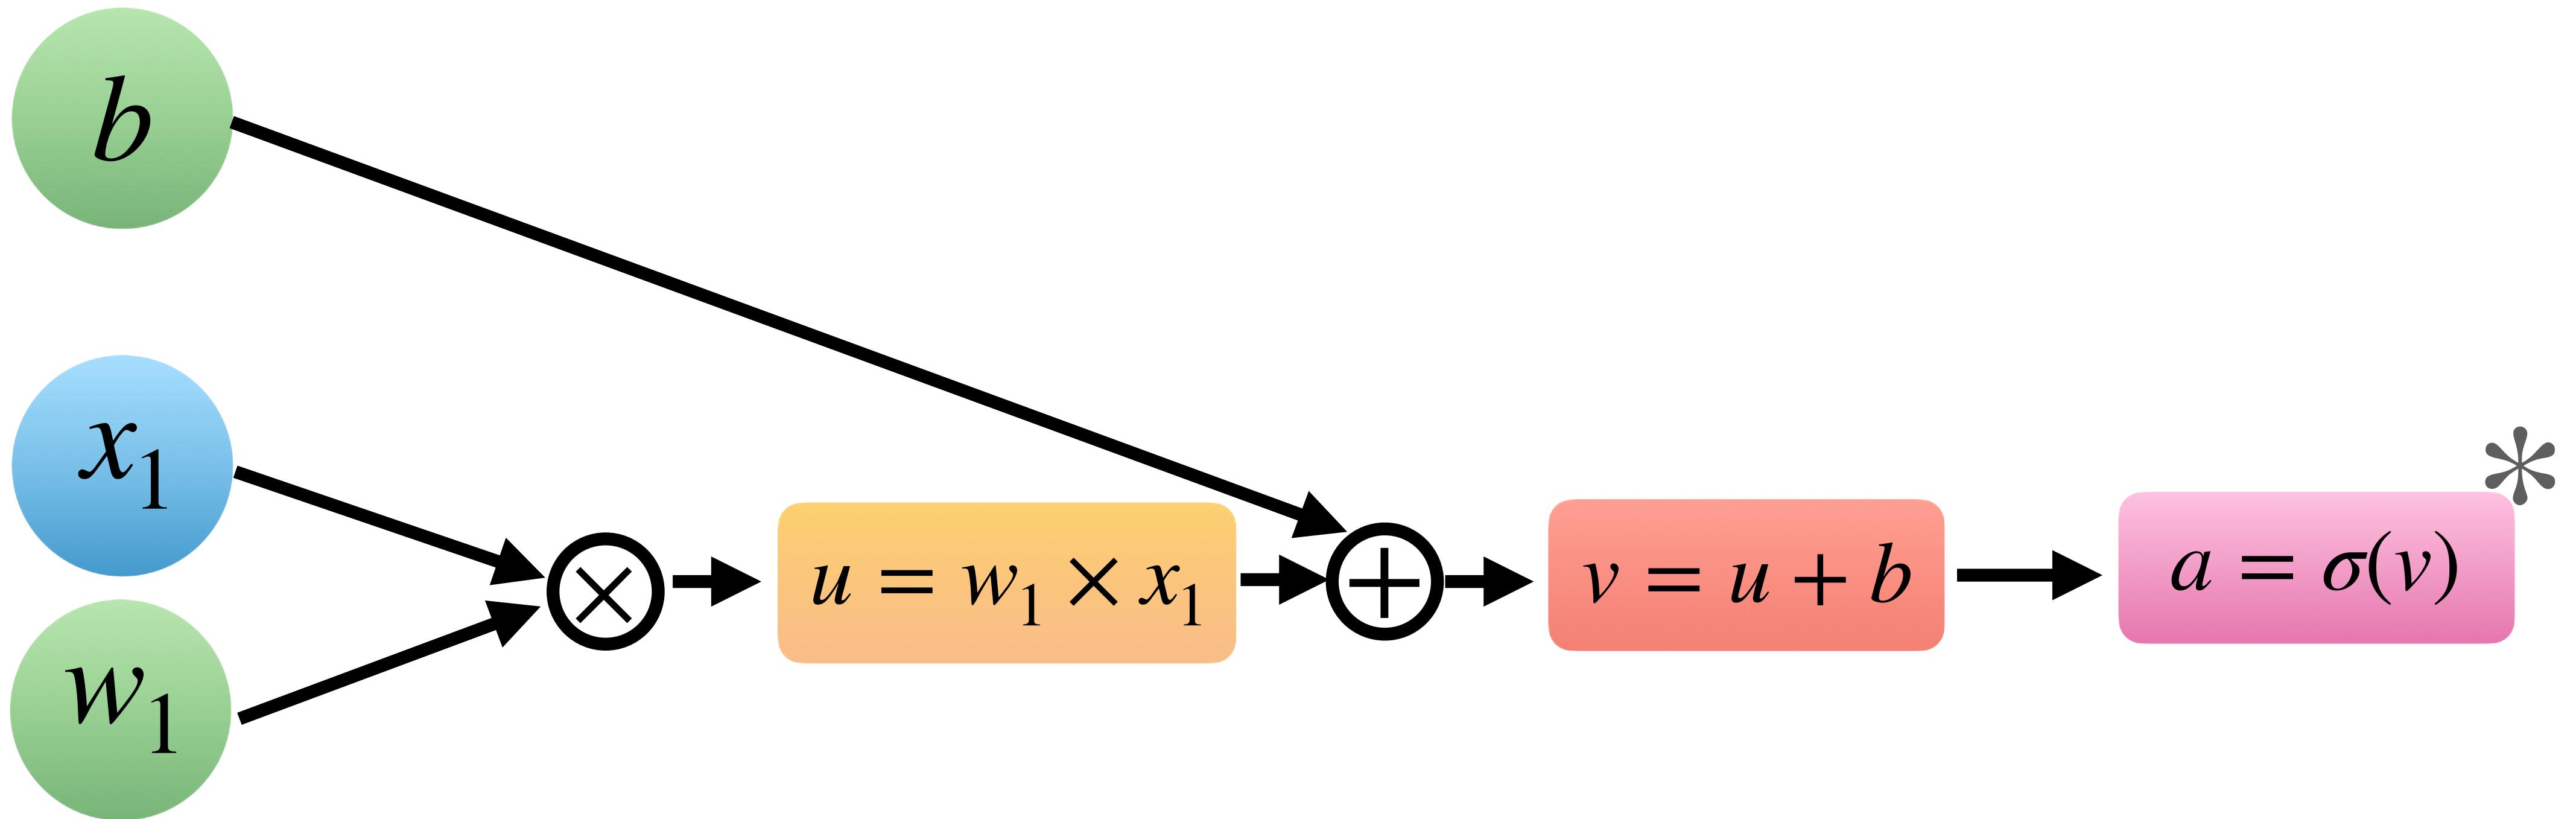
2

Automatic
differentiation engine

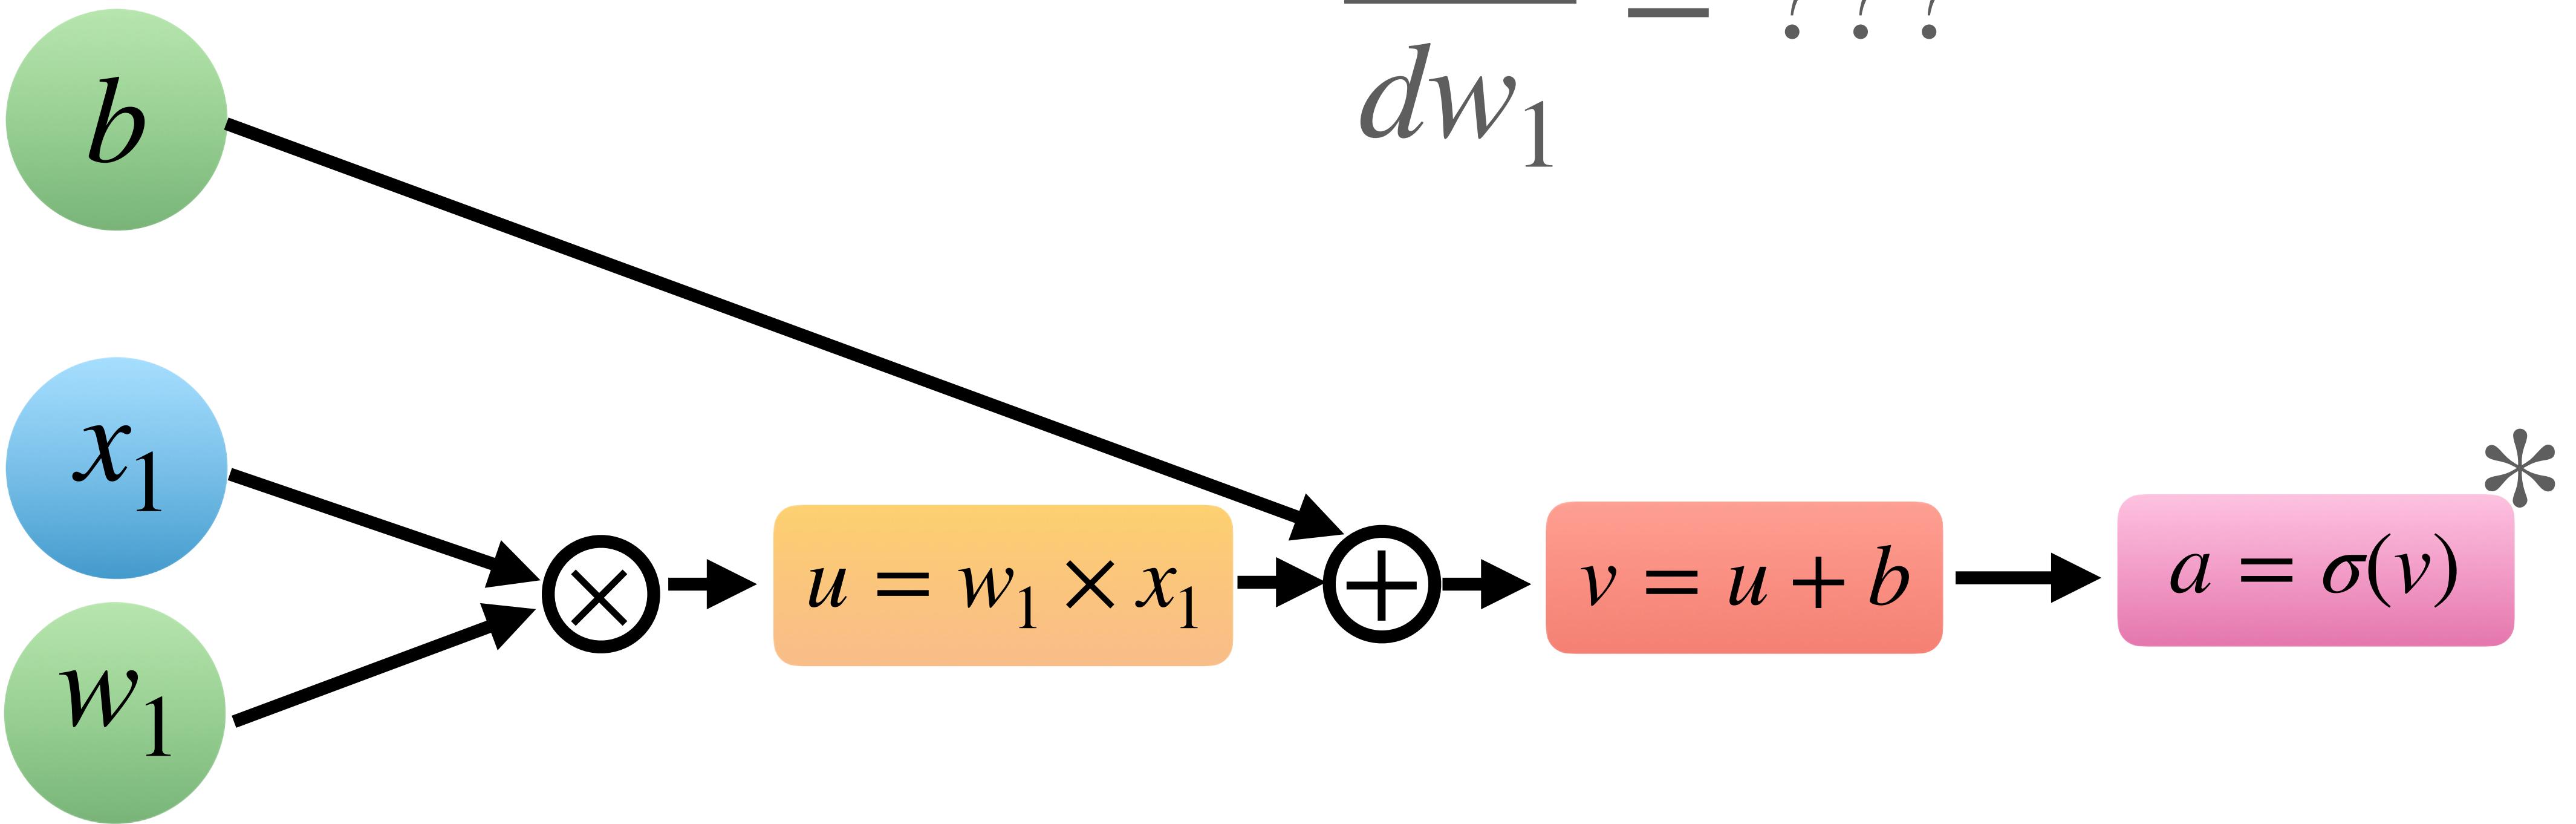
What is PyTorch?





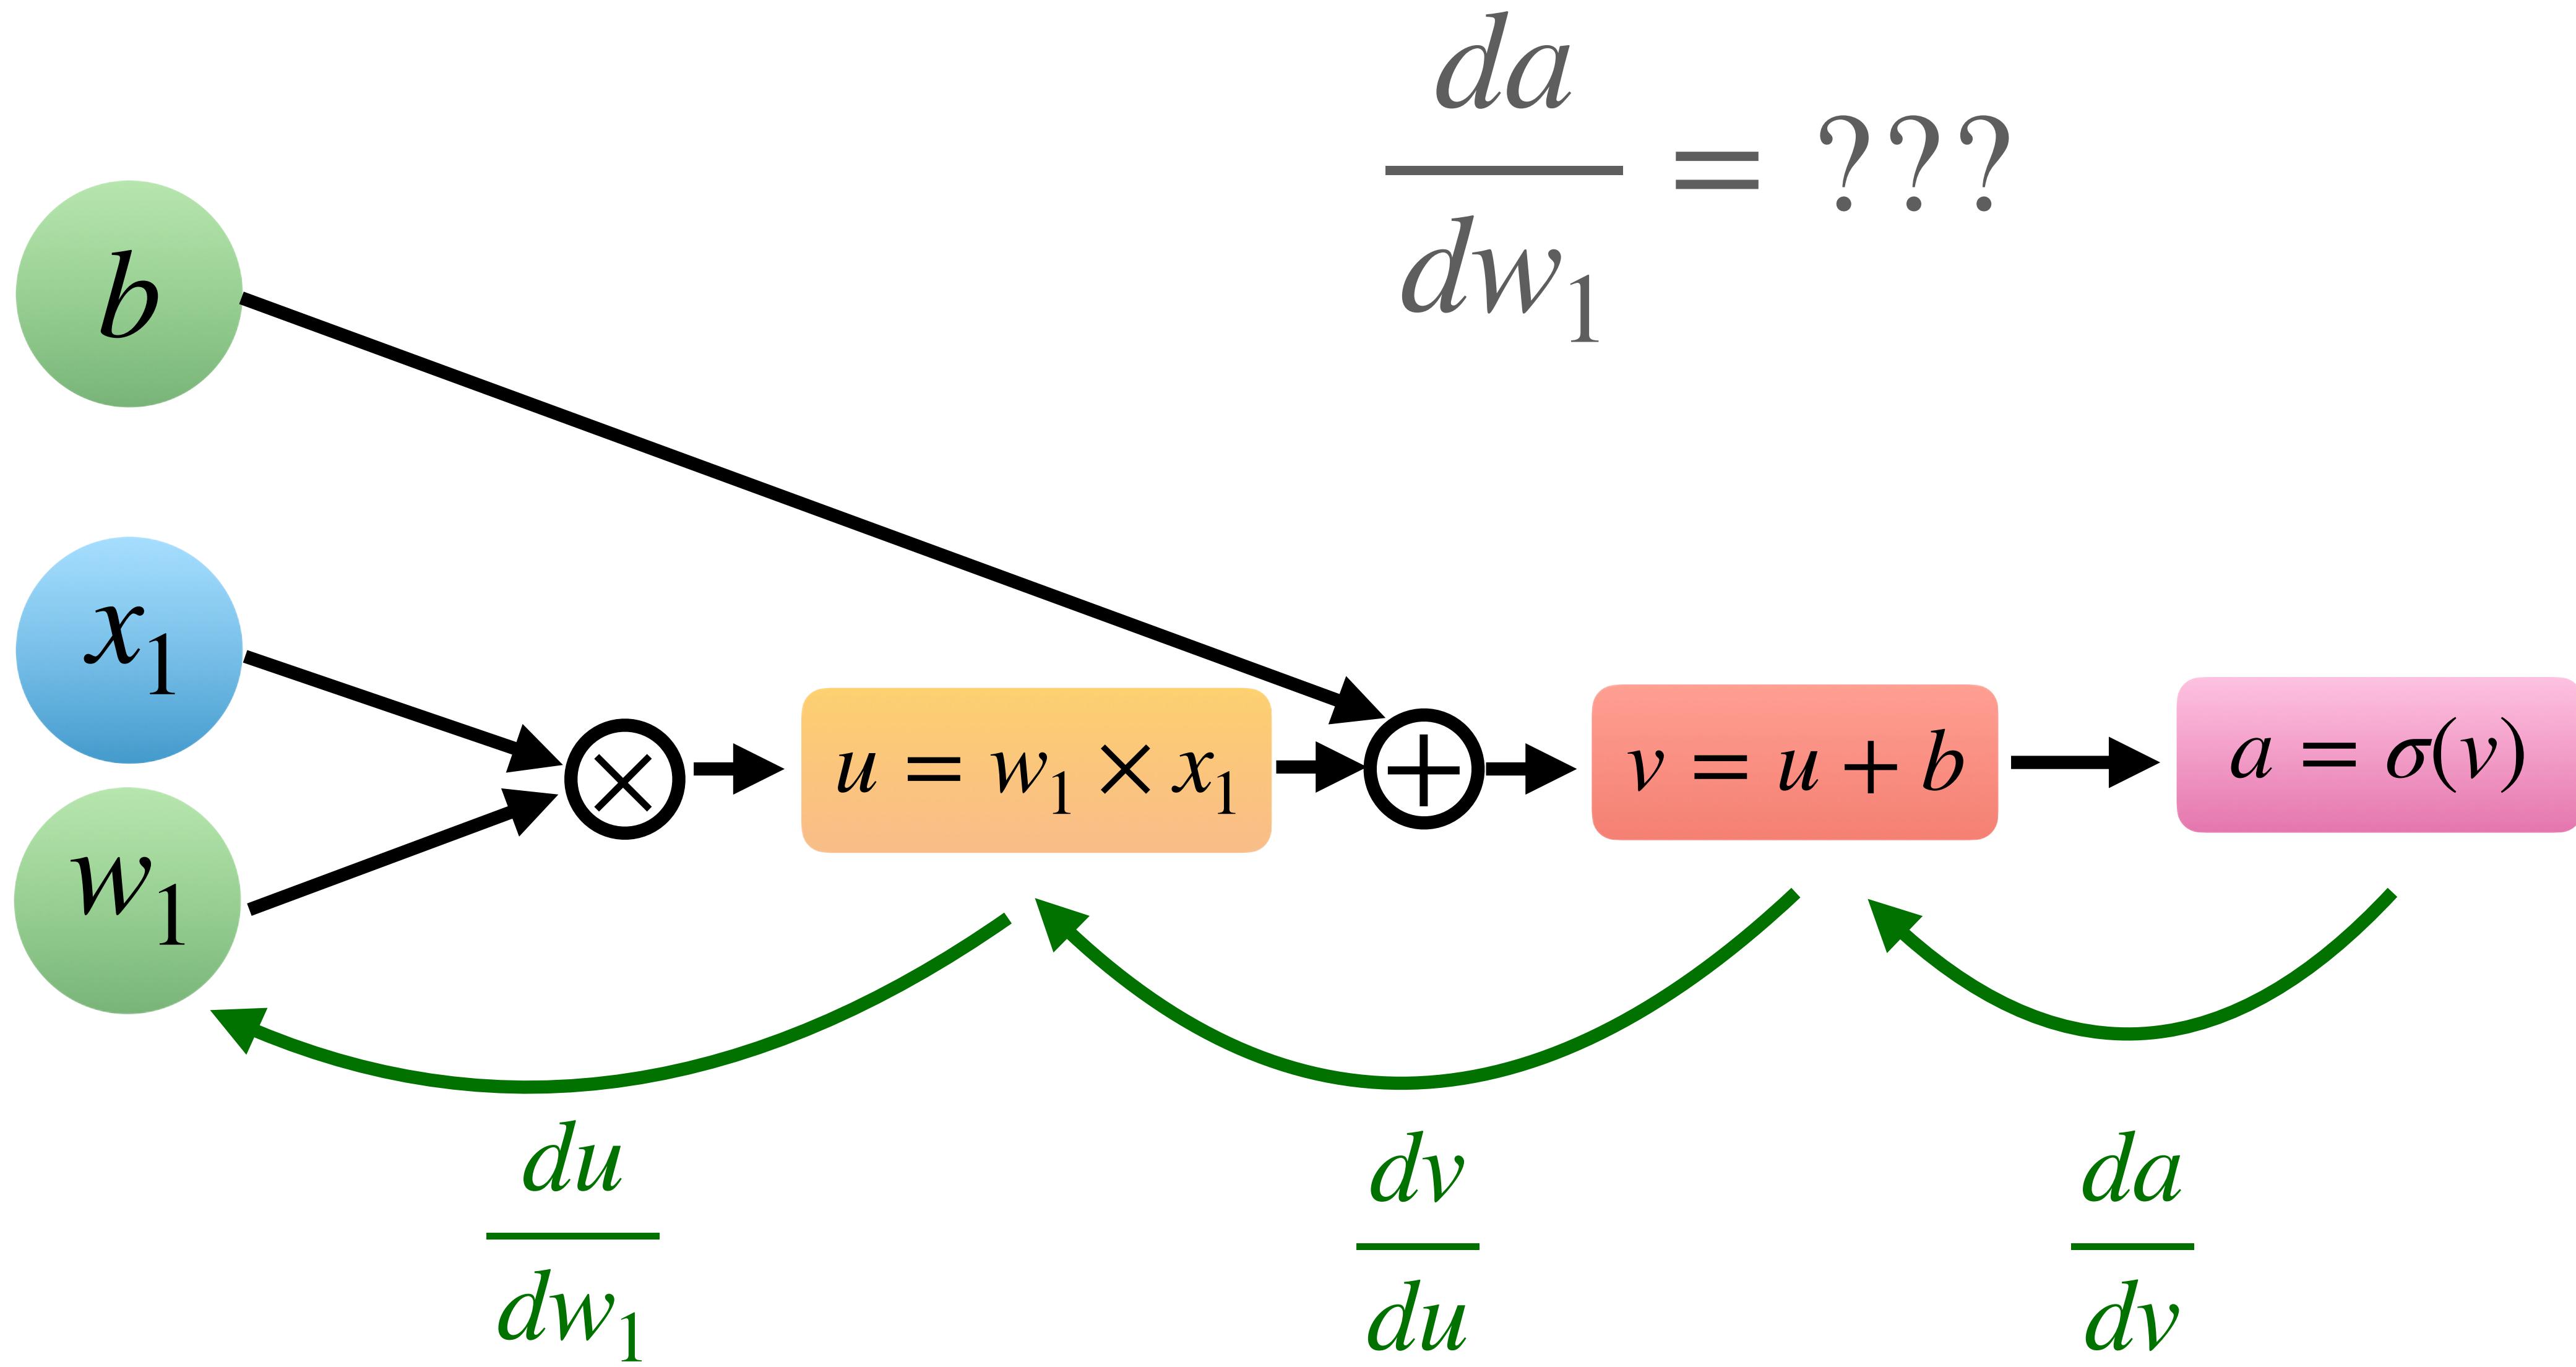


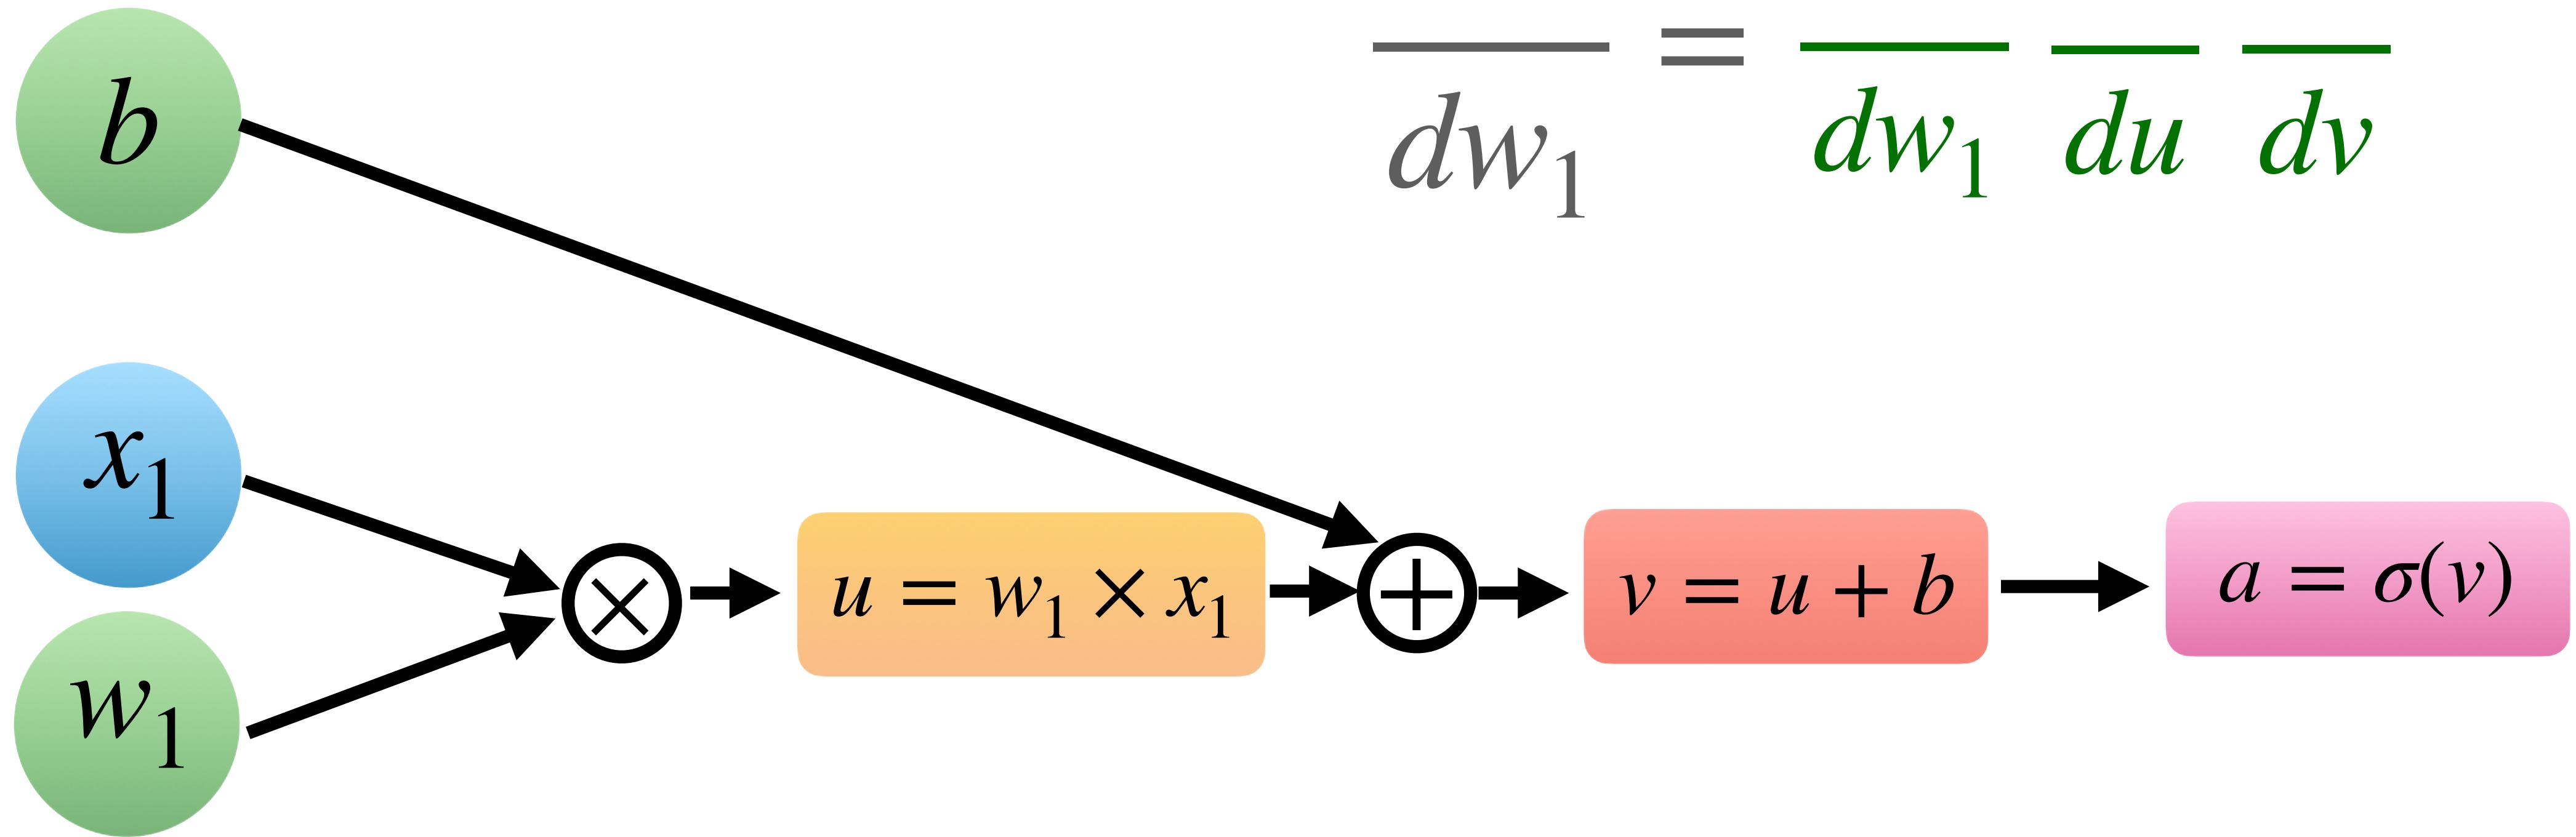
$$*\ \sigma(z) = \frac{1}{1 + e^{-z}}$$

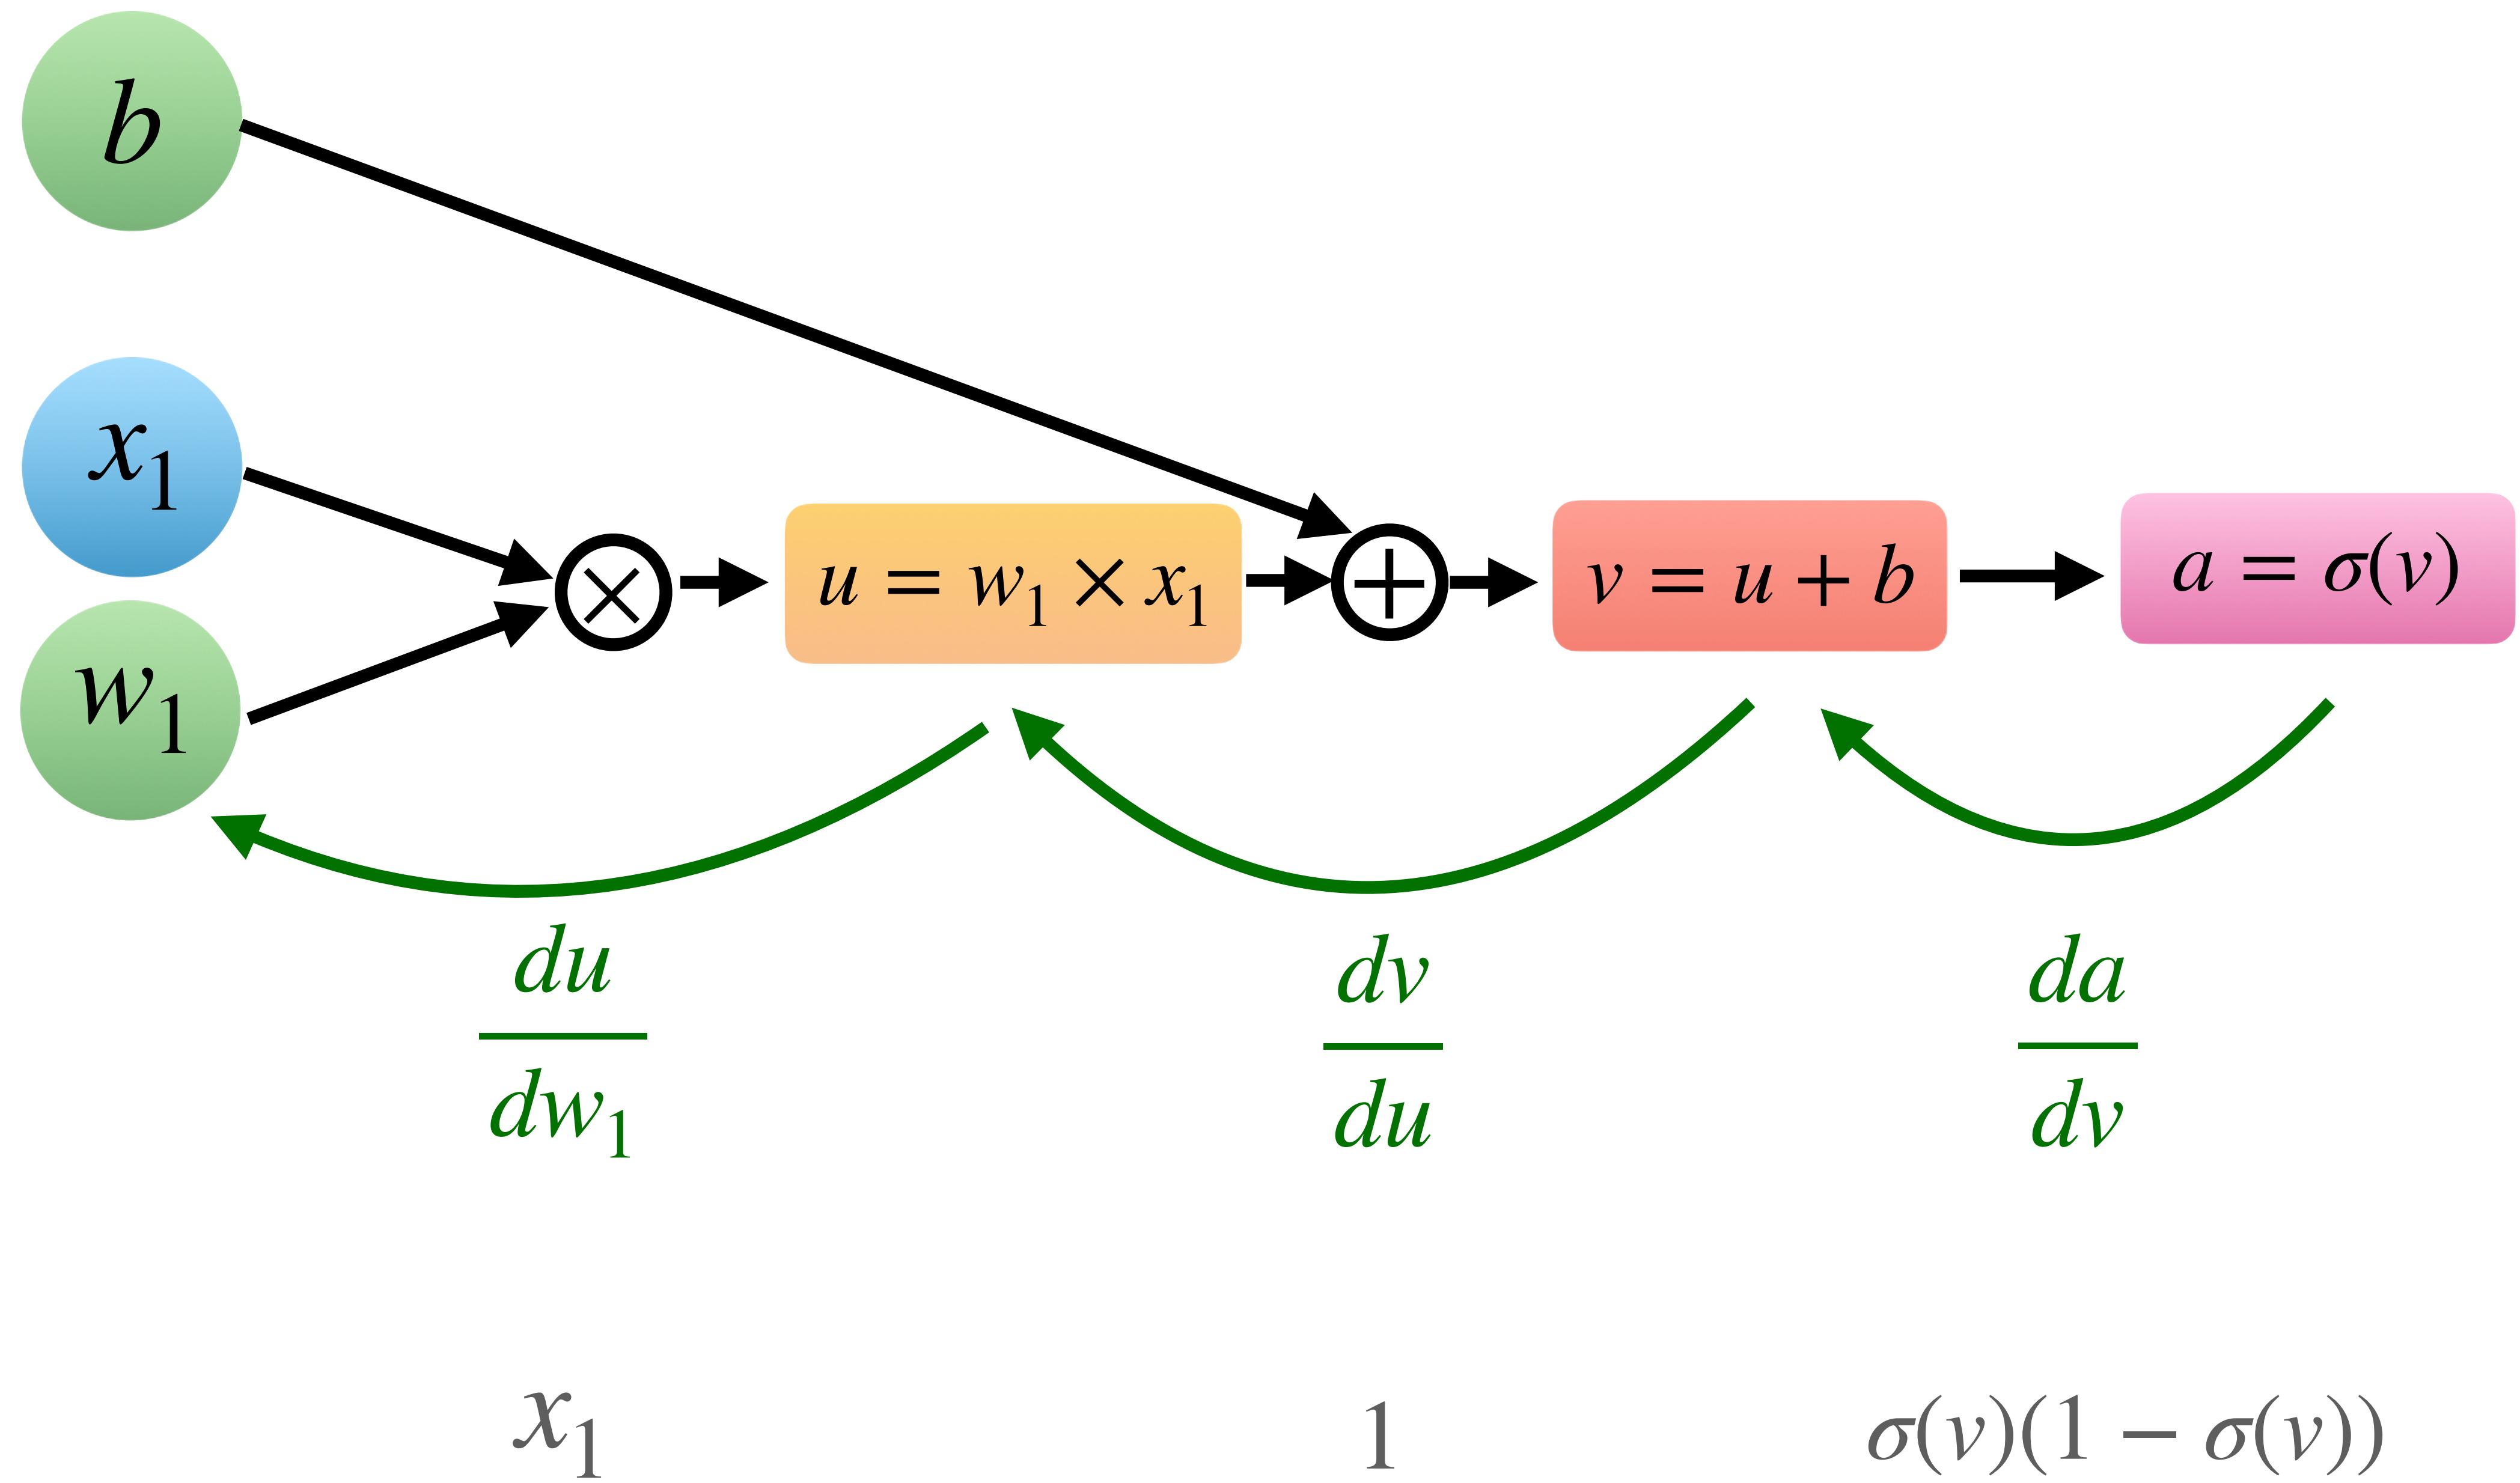


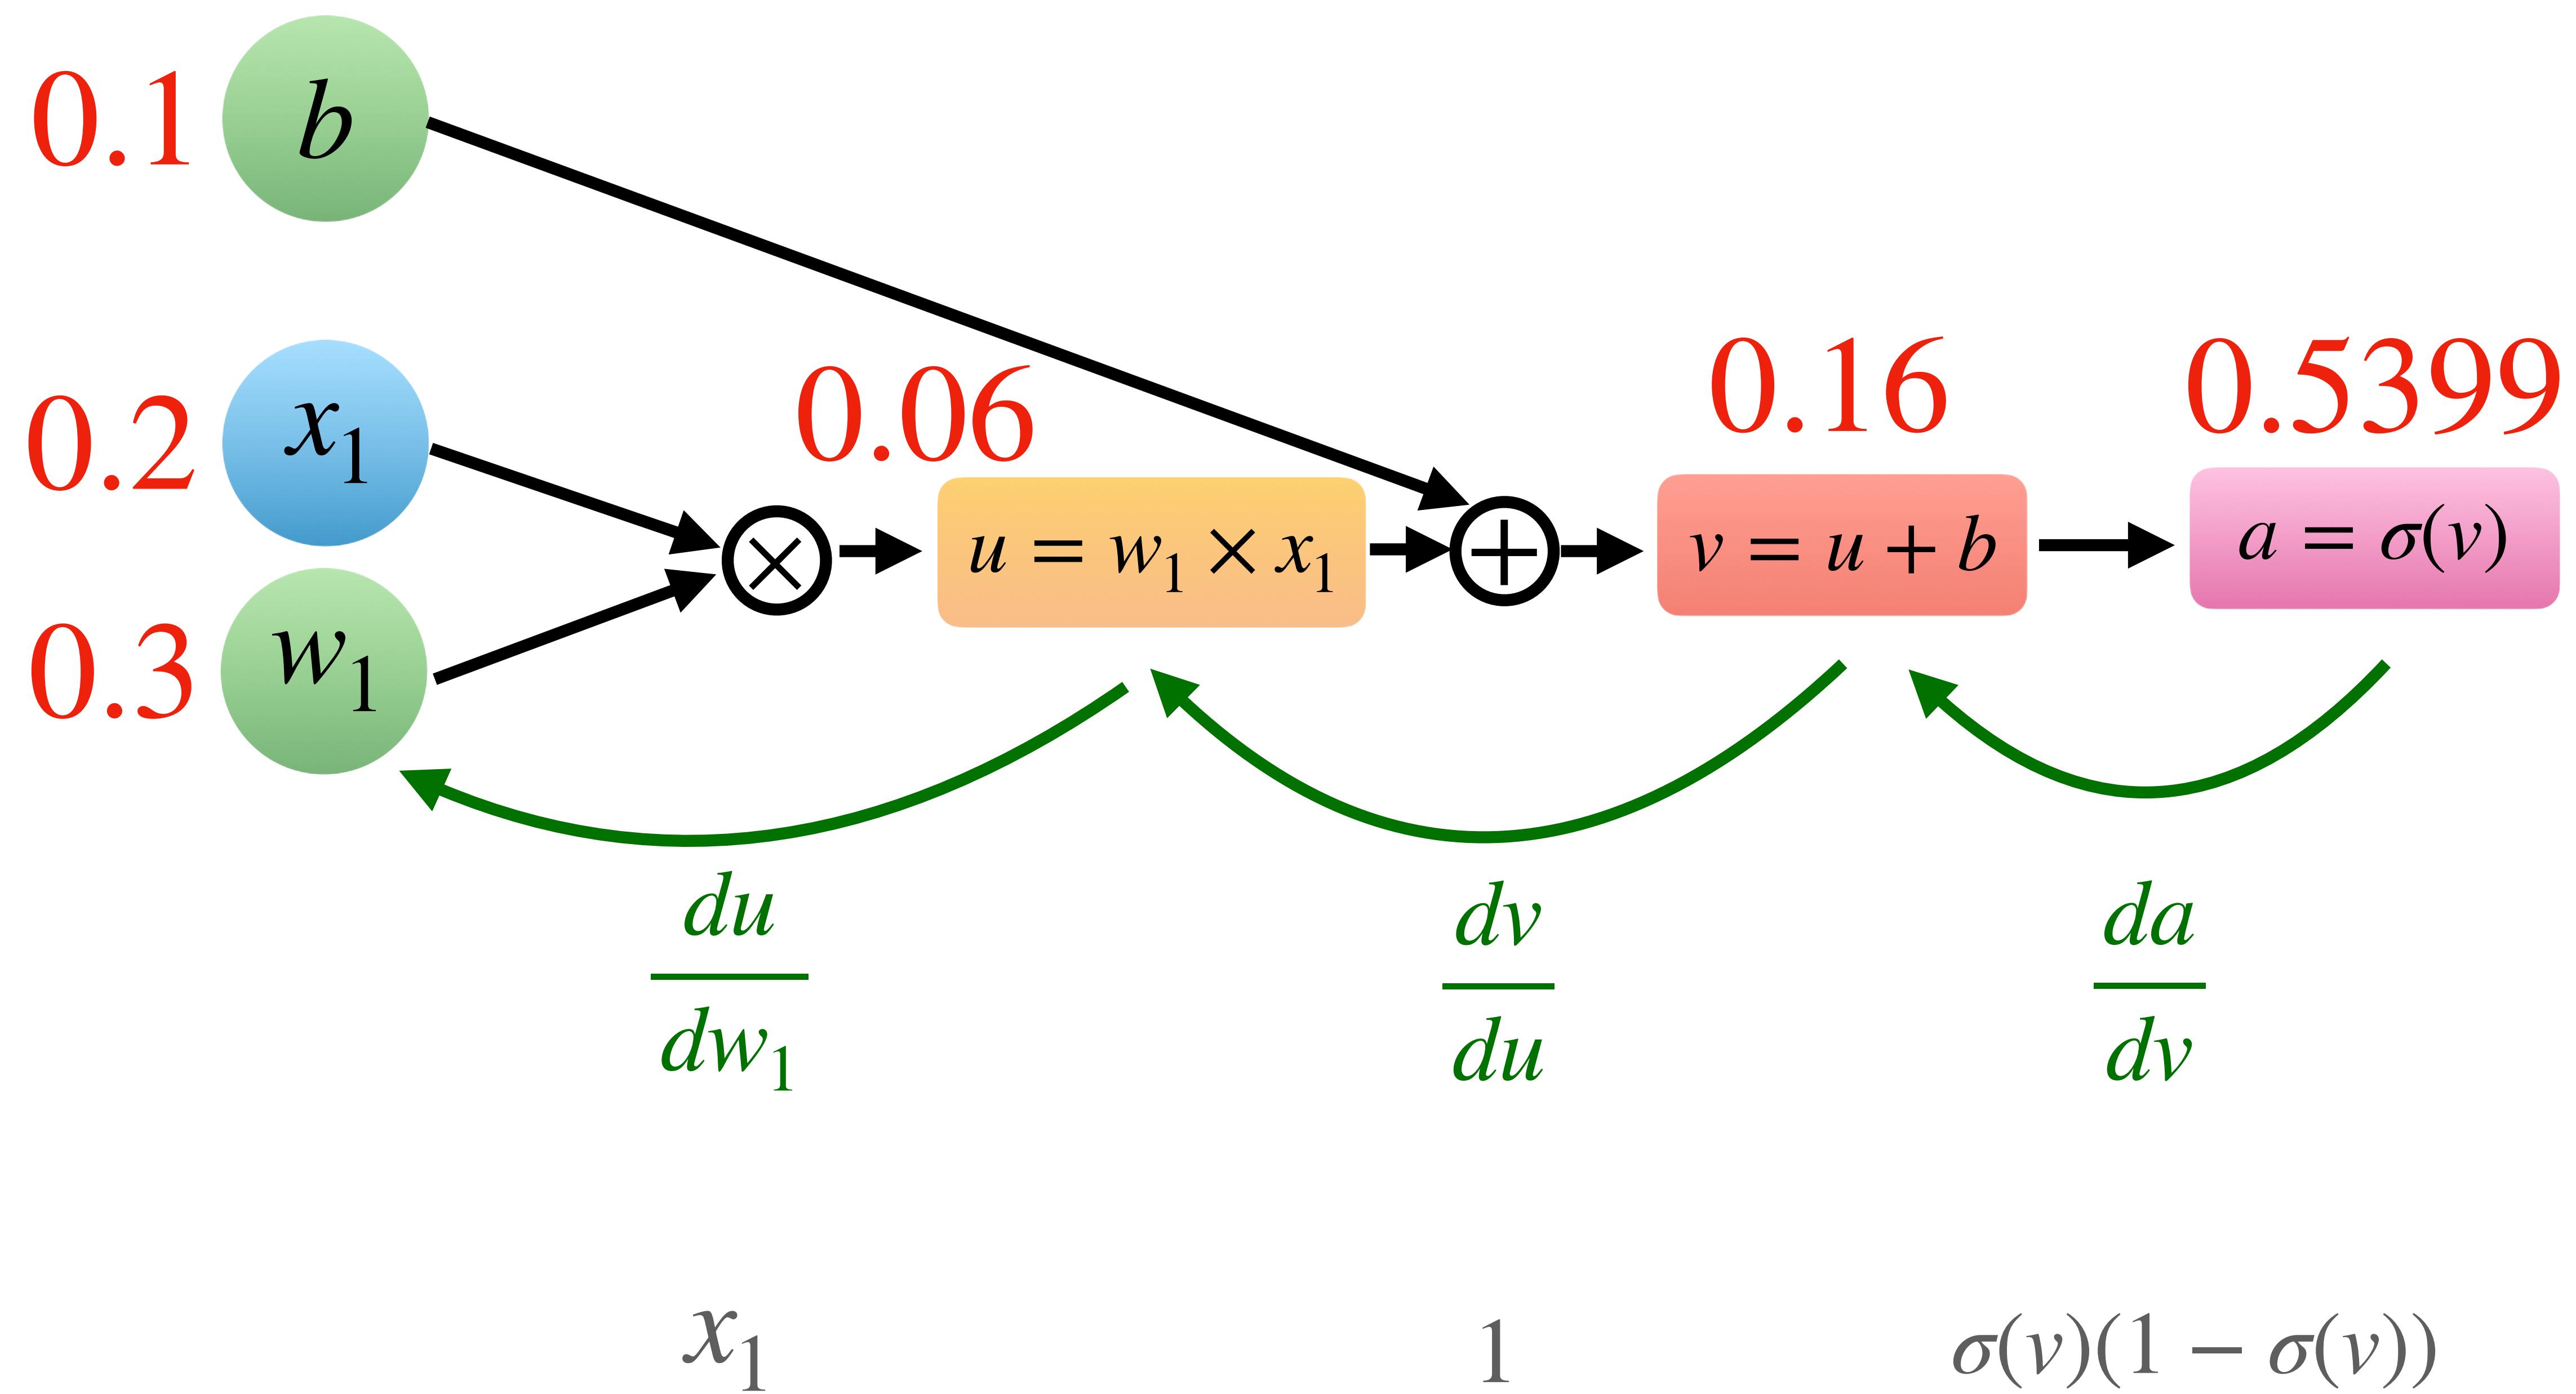
$$\frac{da}{dw_1} = ???$$

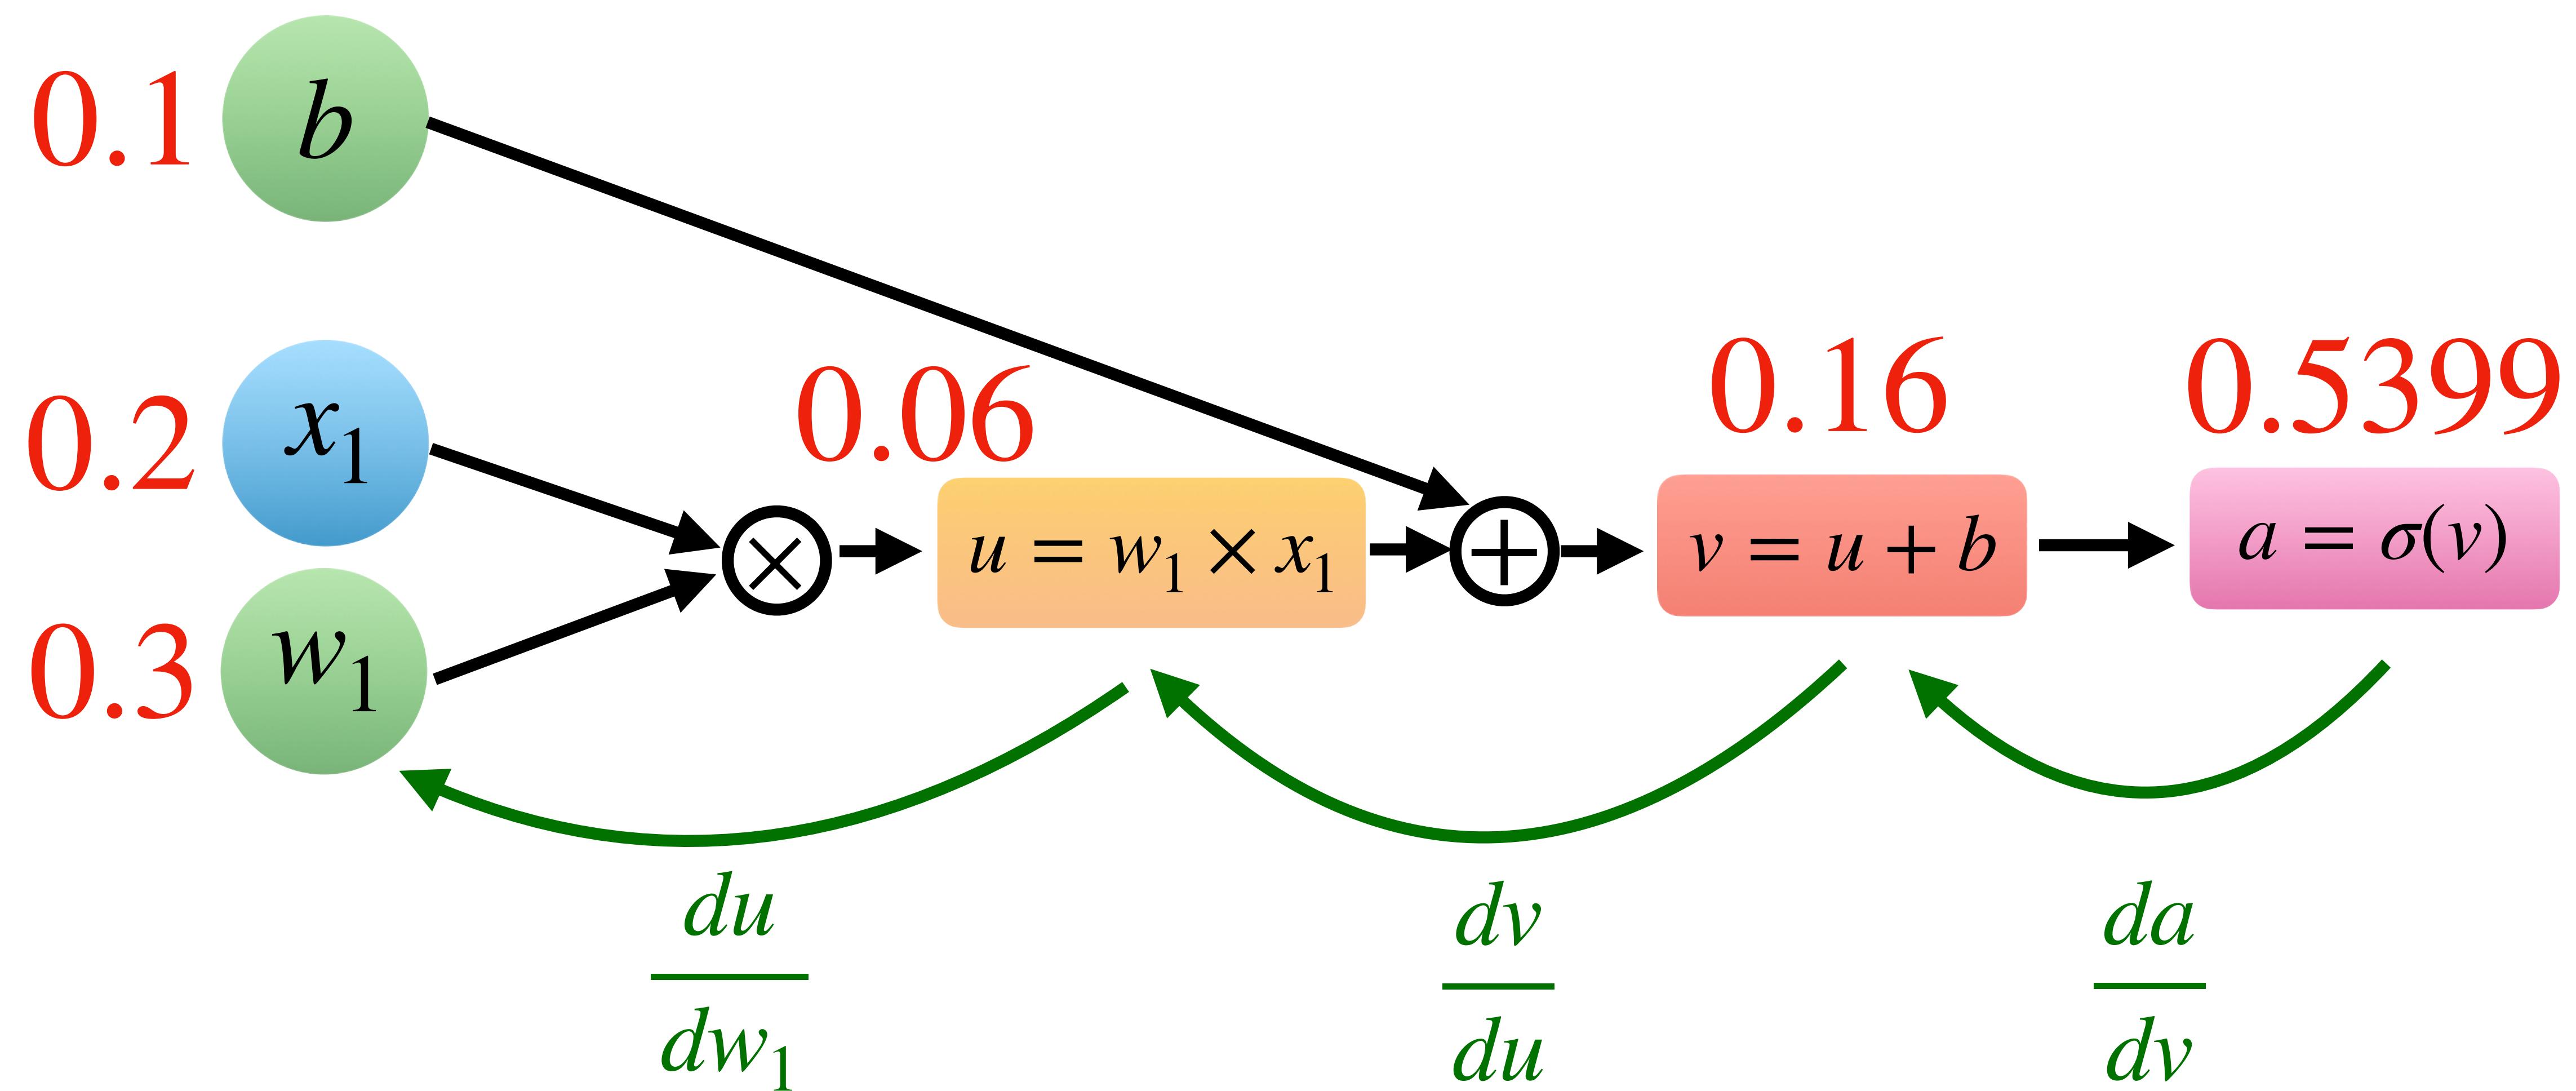
$$*\ \sigma(z) = \frac{1}{1 + e^{-z}}$$





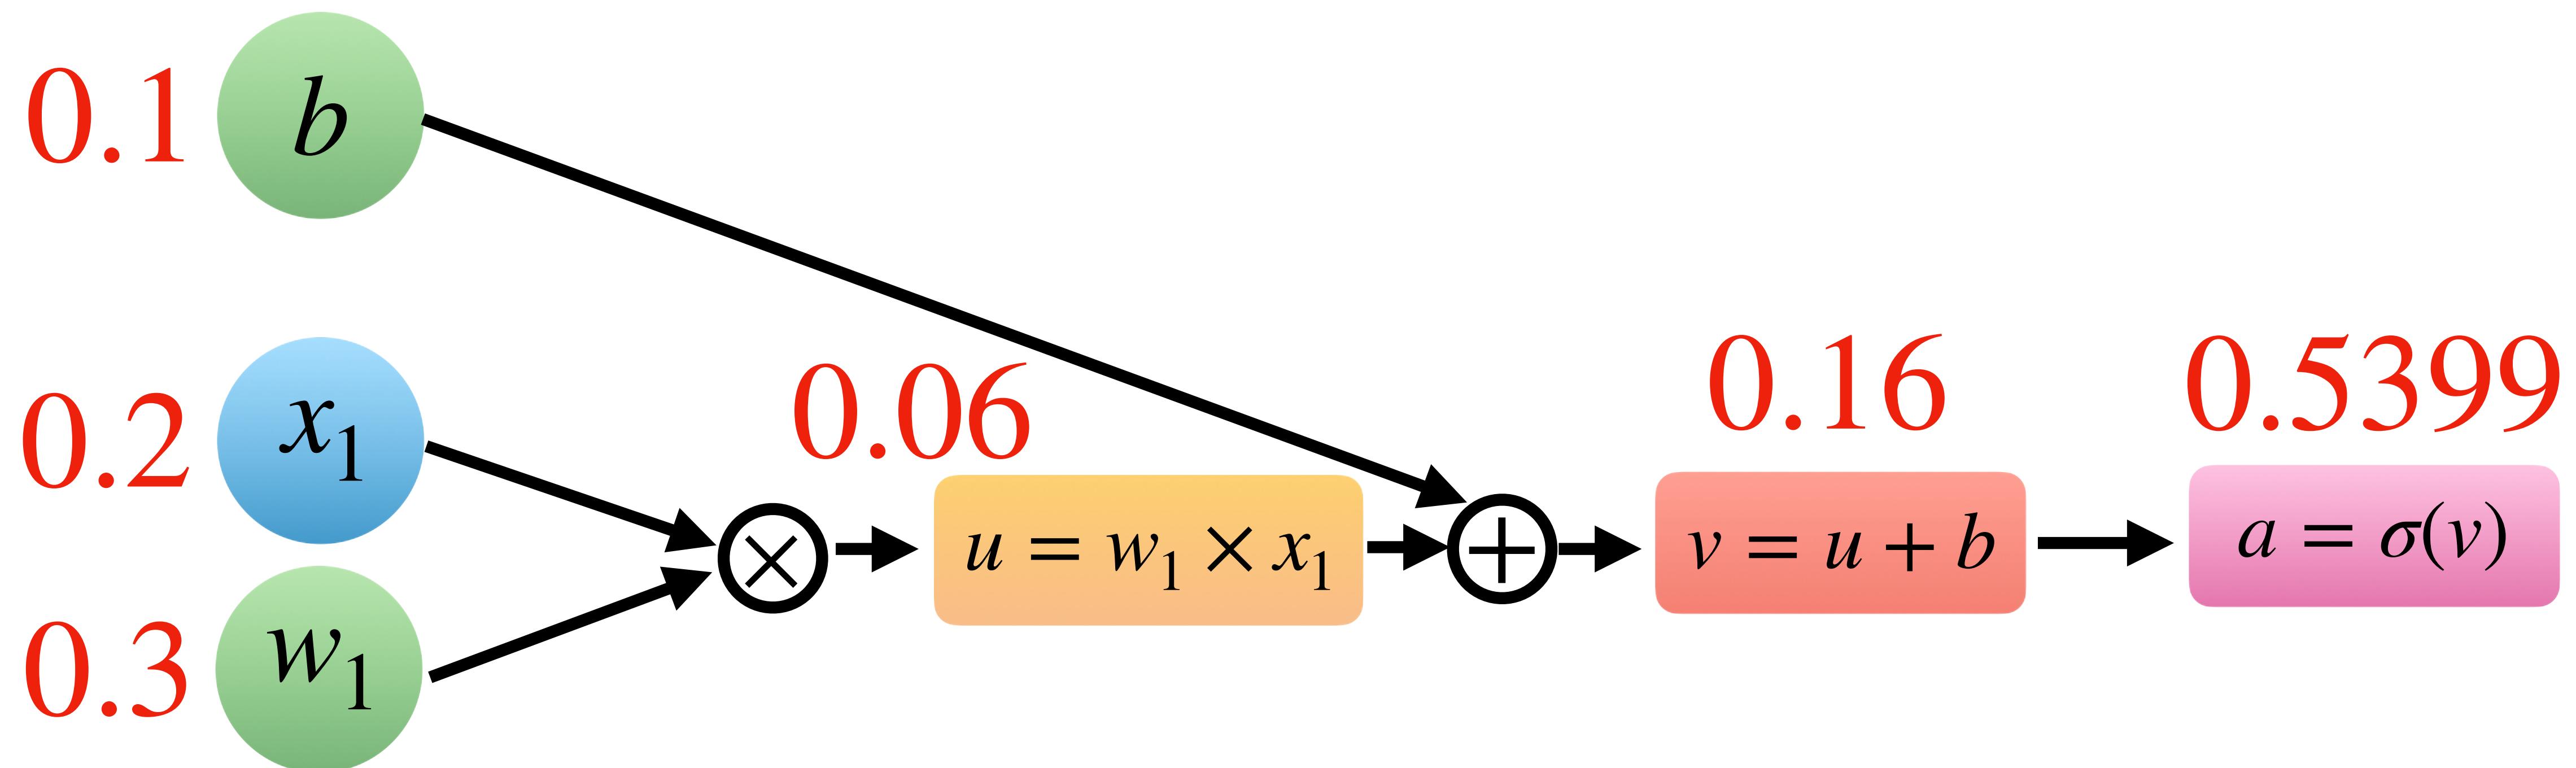






$$\frac{da}{dw_1} = x_1 \times \sigma(v)(1 - \sigma(v))$$

0.0497



Forward pass
in PyTorch

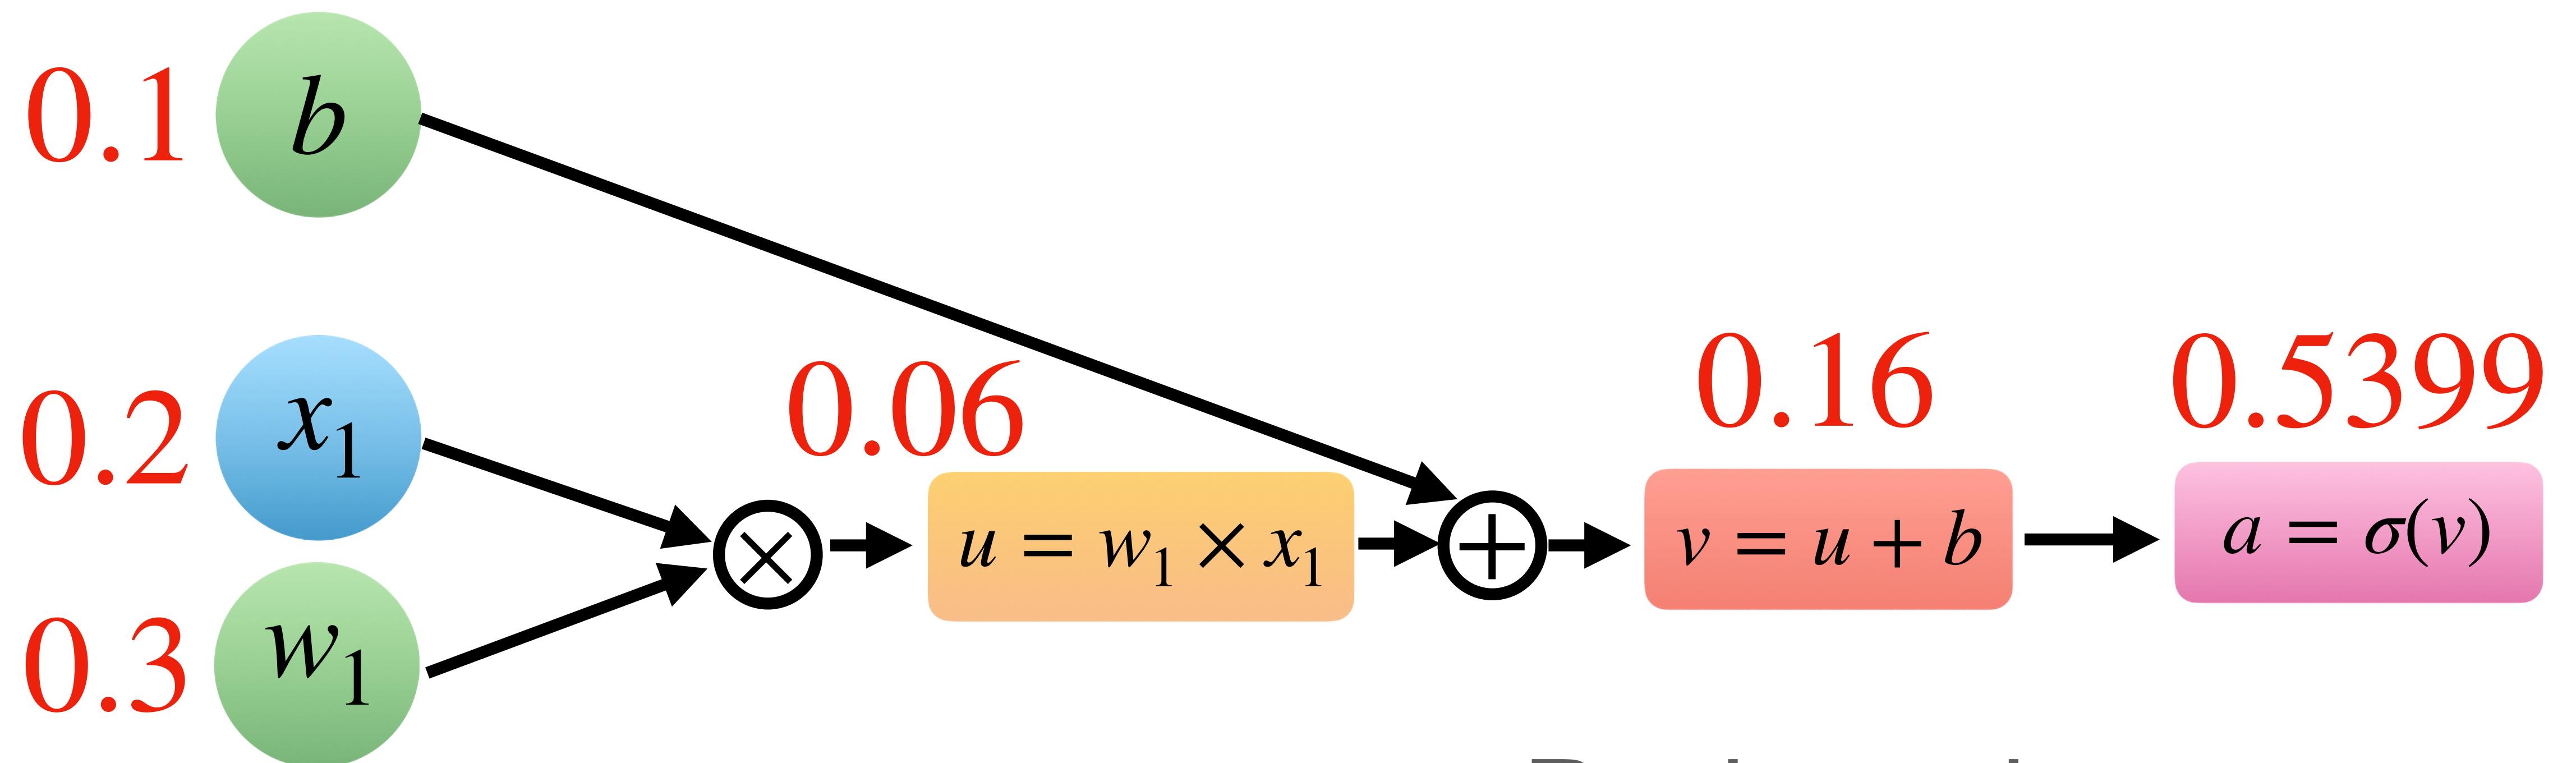
```

b = torch.tensor(0.1)
x1 = torch.tensor(0.2)
w1 = torch.tensor(0.3)

u = w1*x1
v = u + b
a = torch.sigmoid(v)
a

```

`tensor(0.5399)`



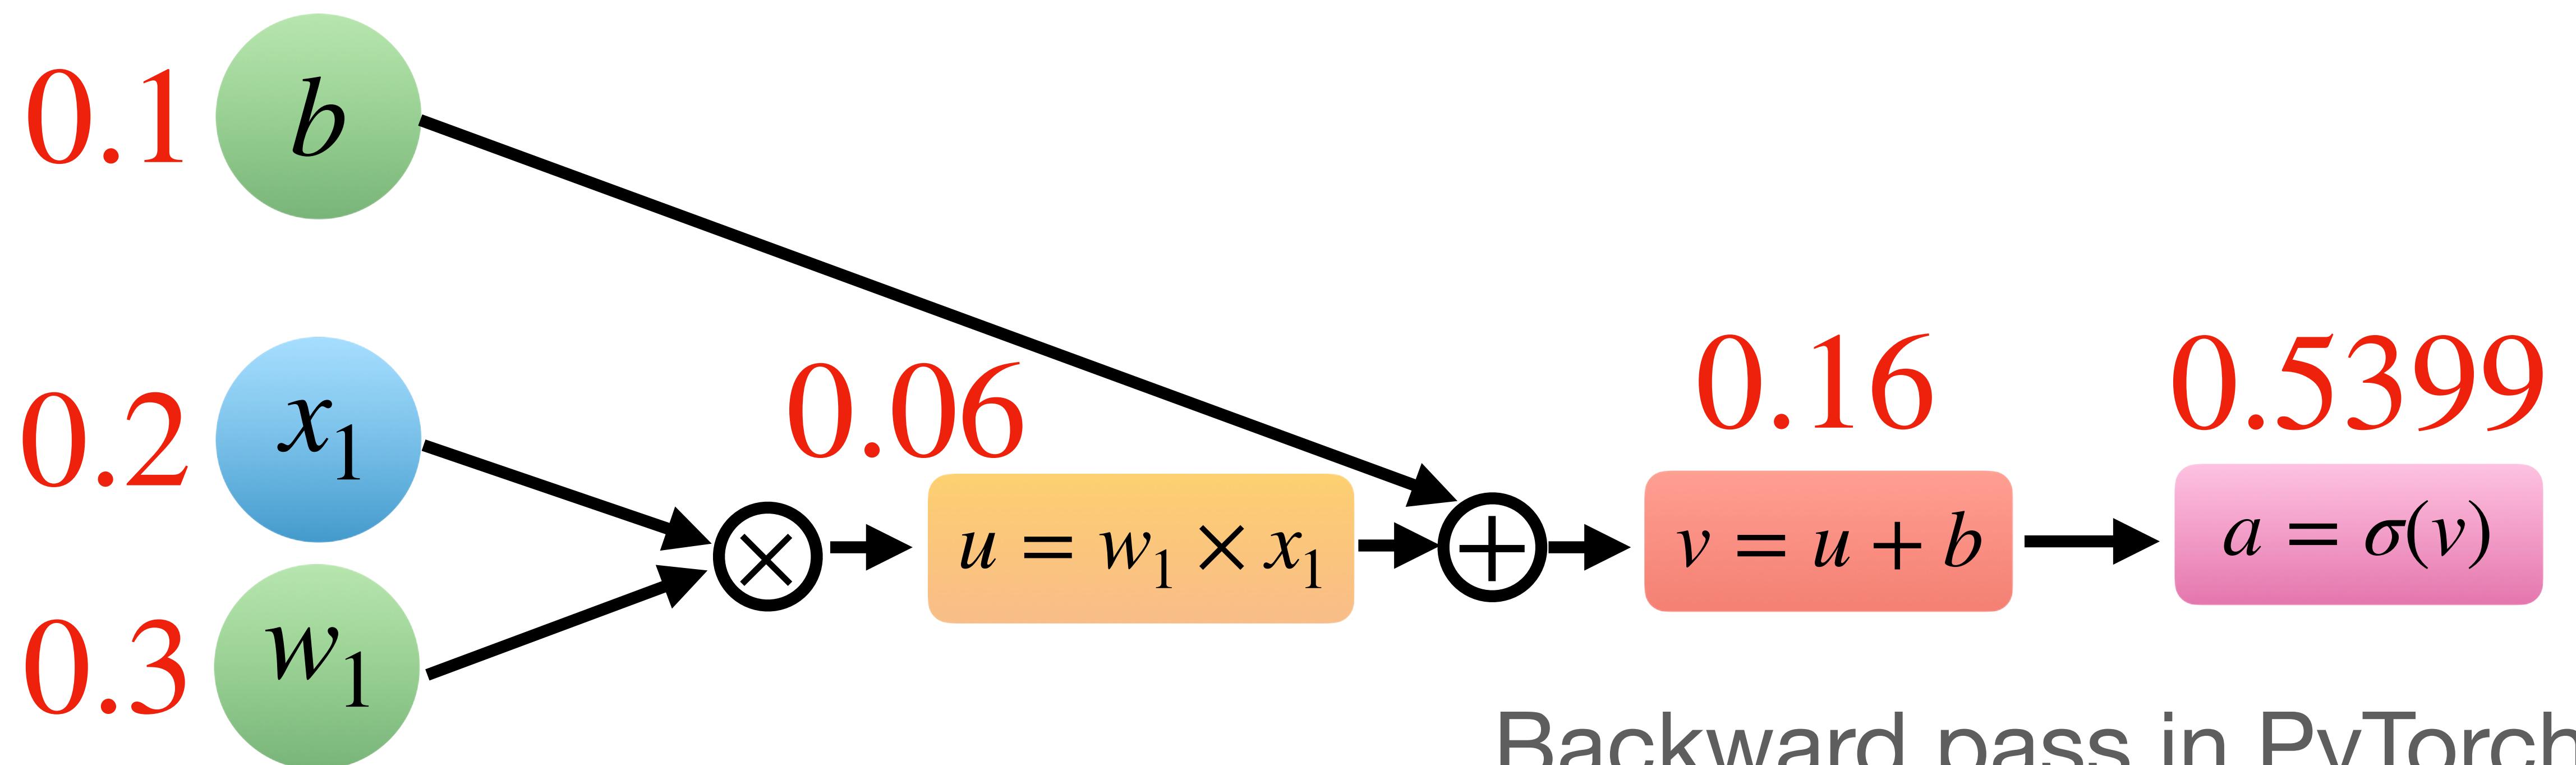
```
b = torch.tensor(0.1)
x1 = torch.tensor(0.2)
w1 = torch.tensor(0.3)
```

```
u = w1*x1
v = u + b
a = torch.sigmoid(v)
a
```

tensor(0.5399)

Backward pass in PyTorch

$$\frac{da}{dw_1} = \begin{aligned} &a * (1-a) * x1 \\ &\text{tensor}(0.0497) \end{aligned}$$



Backward pass in PyTorch
automatically!

```

b = torch.tensor(0.1)
x1 = torch.tensor(0.2)
w1 = torch.tensor(0.3, requires_grad=True)

u = w1*x1
v = u + b
a = torch.sigmoid(v)
a
    
```

`tensor(0.5399, grad_fn=<SigmoidBackward0>)`

$$\frac{da}{dw_1} =$$

```

from torch.autograd import grad
grad(a, w1)
(tensor(0.0497),)
    
```

3

Deep learning
library

What is PyTorch?

PyTorch in 3 steps!

Step 1: Defining the dataset

Step 2: Defining the model

Step 3: Defining the training loop

PyTorch in 3 steps!

Step 1: Defining the model

Step 2: Defining the training loop

Step 3: Defining the dataset

Step 1: Defining the model

Define layers

```
● ● ●  
  
class PyTorchCNN(torch.nn.Module):  
    def __init__(self, num_classes):  
        super().__init__()  
  
        self.conv_1 = torch.nn.Conv2d(...)  
        self.conv_2 = torch.nn.Conv2d(...)  
        self.linear_1 = torch.nn.Linear(..., num_classes)
```

Define layers

Objects with
a **state**
go here

```
● ● ●  
class PyTorchCNN(torch.nn.Module):  
    def __init__(self, num_classes):  
        super().__init__()  
  
        self.conv_1 = torch.nn.Conv2d(...)  
        self.conv_2 = torch.nn.Conv2d(...)  
        self.linear_1 = torch.nn.Linear(..., num_classes)
```

Define forward method



```
class PyTorchCNN(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()

        self.conv_1 = torch.nn.Conv2d(...)
        self.conv_2 = torch.nn.Conv2d(...)
        self.linear_1 = torch.nn.Linear(..., num_classes)

    def forward(self, x):
        out = self.conv_1(x)
        out = torch.nn.functional.relu(out)
        out = torch.nn.functional.max_pool2d(out)

        out = self.conv_2(out)
        out = torch.nn.functional.relu(out)
        out = torch.nn.functional.max_pool2d(out)

        out = torch.flatten(out, dim=1)
        logits = self.linear_1(out)
        return logits
```

Define forward method



```
class PyTorchCNN(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()

        self.conv_1 = torch.nn.Conv2d(...)
        self.conv_2 = torch.nn.Conv2d(...)
        self.linear_1 = torch.nn.Linear(..., num_classes)
```

```
def forward(self, x):
    out = self.conv_1(x)
    out = torch.nn.functional.relu(out)
    out = torch.nn.functional.max_pool2d(out)

    out = self.conv_2(out)
    out = torch.nn.functional.relu(out)
    out = torch.nn.functional.max_pool2d(out)

    out = torch.flatten(out, dim=1)
    logits = self.linear_1(out)
    return logits
```

How things
are executed

Define forward method



```
class PyTorchCNN(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()

        self.conv_1 = torch.nn.Conv2d(...)
        self.conv_2 = torch.nn.Conv2d(...)
        self.linear_1 = torch.nn.Linear(..., num_classes)

    def forward(self, x):
        out = self.conv_1(x)
        out = torch.nn.functional.relu(out)
        out = torch.nn.functional.max_pool2d(out)

        out = self.conv_2(out)
        out = torch.nn.functional.relu(out)
        out = torch.nn.functional.max_pool2d(out)

        out = torch.flatten(out, dim=1)
        logits = self.linear_1(out)
        return logits
```

How things
are executed

Functional
API since
these don't
need a **state**



```
class PyTorchCNN(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()

        self.conv_1 = torch.nn.Conv2d(...)
        self.conv_2 = torch.nn.Conv2d(...)
        self.linear_1 = torch.nn.Linear(..., num_classes)

    def forward(self, x):
        out = self.conv_1(x)
        out = torch.nn.functional.relu(out)
        out = torch.nn.functional.max_pool2d(out)

        out = self.conv_2(out)
        out = torch.nn.functional.relu(out)
        out = torch.nn.functional.max_pool2d(out)

        out = torch.flatten(out, dim=1)
        logits = self.linear_1(out)
        return logits
```

Initialize
the model

```
torch.manual_seed(random_seed)
model = PyTorchCNN(num_classes=num_classes)
model = model.to(device)
```

Make your life easier with Sequential()

Make your life
easier with
`Sequential()`

```
class PyTorchCNN(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()

        self.num_classes = num_classes
        self.features = torch.nn.Sequential(
            torch.nn.Conv2d(...),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(...),
            torch.nn.Conv2d(...),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(...),
        )
        self.classifier = torch.nn.Sequential(
            torch.nn.Flatten(),
            torch.nn.Linear(..., num_classes)
        )

    def forward(self, x):
        x = self.features(x)
        x = self.classifier(x)
        return x
```

Definition order
equals
execution order

```
class PyTorchCNN(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()

        self.num_classes = num_classes
        self.features = torch.nn.Sequential(
            torch.nn.Conv2d(...),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(...),
            torch.nn.Conv2d(...),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(...),
        )
        self.classifier = torch.nn.Sequential(
            torch.nn.Flatten(),
            torch.nn.Linear(..., num_classes)
        )

    def forward(self, x):
        x = self.features(x)
        x = self.classifier(x)
        return x
```

Definition order
equals
execution order

```
class PyTorchCNN(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()

        self.num_classes = num_classes
        self.features = torch.nn.Sequential(
            torch.nn.Conv2d(...),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(...),
            torch.nn.Conv2d(...),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(...),
        )
        self.classifier = torch.nn.Sequential(
            torch.nn.Flatten(),
            torch.nn.Linear(..., num_classes)
        )

    def forward(self, x):
        x = self.features(x)
        x = self.classifier(x)
        return x
```

Step 2: Defining the training loop

Initializing the model and optimizer



```
model = PyTorchCNN(num_classes=num_classes)
model = model.to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```

Iterating over the training examples



```
model = PyTorchCNN(num_classes=num_classes)
model = model.to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epoch in range(num_epochs):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features, targets = features.to(device), targets.to(device)
```

Computing the predictions

```
● ● ●

model = PyTorchCNN(num_classes=num_classes)
model = model.to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epoch in range(num_epochs):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features, targets = features.to(device), targets.to(device)

        ### Forward pass
        logits = model(features)
        loss = F.cross_entropy(logits, targets)
```

Computing the predictions

```
● ● ●  
  
model = PyTorchCNN(num_classes=num_classes)  
model = model.to(device)  
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)  
  
for epoch in range(num_epochs):  
    model.train()  
    for batch_idx, (features, targets) in enumerate(train_loader):  
  
        features, targets = features.to(device), targets.to(device)  
  
        ### Forward pass  
        logits = model(features)  
        loss = F.cross_entropy(logits, targets)
```

the `cross_entropy` loss takes care of the LogSoftmax internally

Computing the predictions

```
● ● ●  
  
model = PyTorchCNN(num_classes=num_classes)  
model = model.to(device)  
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)  
  
for epoch in range(num_epochs):  
    model.train()  
    for batch_idx, (features, targets) in enumerate(train_loader):  
  
        features, targets = features.to(device), targets.to(device)  
  
        ### Forward pass  
        logits = model(features)  
        loss = F.cross_entropy(logits, targets)
```

the `cross_entropy` loss also takes care of one-hot encoding internally

Computing the backward pass



```
model = PyTorchCNN(num_classes=num_classes)
model = model.to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epoch in range(num_epochs):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features, targets = features.to(device), targets.to(device)

        ### Forward pass
        logits = model(features)
        loss = F.cross_entropy(logits, targets)

        ### Backward pass (backpropagation)
        optimizer.zero_grad()
        loss.backward()
```

Computing the backward pass



```
model = PyTorchCNN(num_classes=num_classes)
model = model.to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epoch in range(num_epochs):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features, targets = features.to(device), targets.to(device)

        ### Forward pass
        logits = model(features)
        loss = F.cross_entropy(logits, targets)

        ### Backward pass (backpropagation)
        optimizer.zero_grad()
        loss.backward()
```

If you omit this, gradients will be accumulated, which we usually don't want

Updating the model weights



```
model = PyTorchCNN(num_classes=num_classes)
model = model.to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epoch in range(num_epochs):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features, targets = features.to(device), targets.to(device)

        ### Forward pass
        logits = model(features)
        loss = F.cross_entropy(logits, targets)

        ### Backward pass (backpropagation)
        optimizer.zero_grad()
        loss.backward()

        ### Update model parameters
        optimizer.step()
```

Tracking the perf- mance

```
model = PyTorchCNN(num_classes=num_classes)
model = model.to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epoch in range(num_epochs):
    model.train()
    for batch_idx, (features, targets) in enumerate(train_loader):

        features, targets = features.to(device), targets.to(device)

        ### Forward pass
        logits = model(features)
        loss = F.cross_entropy(logits, targets)

        ### Backward pass (backpropagation)
        optimizer.zero_grad()
        loss.backward()

        ### Update model parameters
        optimizer.step()

    ### Optional evaluation steps
    model.eval()
    with torch.no_grad():
        valid_acc = compute_accuracy(model, valid_loader, device)
        print(f"Validation accuracy: {valid_acc * 100:.2f}%")
```

Step 3: Defining the dataset

MNIST: The “Hello World” of deep learning



MNIST:

The “Hello World” of deep learning



▼	test	Today at 10:51 AM	-- Folder
>	0	Today at 10:51 AM	-- Folder
>	1	Today at 10:51 AM	-- Folder
>	2	Today at 10:51 AM	-- Folder
>	3	Today at 10:51 AM	-- Folder
>	4	Today at 10:51 AM	-- Folder
>	5	Today at 10:51 AM	-- Folder
>	6	Today at 10:51 AM	-- Folder
>	7	Today at 10:51 AM	-- Folder
>	8	Today at 10:51 AM	-- Folder
>	9	Today at 10:51 AM	-- Folder
▼	train	Today at 10:51 AM	-- Folder
>	0	Today at 10:51 AM	-- Folder
>	1	Today at 10:51 AM	-- Folder
>	2	Today at 10:51 AM	-- Folder
>	3	Today at 10:51 AM	-- Folder
>	4	Today at 10:51 AM	-- Folder
>	5	Today at 10:51 AM	-- Folder
>	6	Today at 10:51 AM	-- Folder
>	7	Today at 10:51 AM	-- Folder
>	8	Today at 10:51 AM	-- Folder
>	9	Today at 10:51 AM	-- Folder

Building the dataset from folder structures



```
from torch.utils.data.dataset import random_split
from torchvision.datasets import ImageFolder

train_dset = ImageFolder(root="mnist-pngs/train", transform=data_transforms["train"])

train_dset, valid_dset = random_split(train_dset, lengths=[55000, 5000])

test_dset = ImageFolder(root="mnist-pngs/test", transform=data_transforms["test"])
```

Defining optional transformations

```
from torch.utils.data.dataset import random_split
from torchvision.datasets import ImageFolder

train_dset = ImageFolder(root="mnist-pngs/train", transform=data_transforms["train"])

train_dset, valid_dset = random_split(train_dset, lengths=[55000, 5000])

test_dset = ImageFolder(root="mnist-pngs/test", transform=data_transforms["test"])
```

```
from torchvision import transforms

data_transforms = {
    "train": transforms.Compose(
        [
            transforms.Resize(32),
            transforms.RandomCrop((28, 28)),
            transforms.ToTensor(),
            # normalize images to [-1, 1] range
            transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
        ]
    ),
}
```

Defining optional transformations

```
from torch.utils.data.dataset import random_split
from torchvision.datasets import ImageFolder

train_dset = ImageFolder(root="mnist-pngs/train", transform=data_transforms["train"])

train_dset, valid_dset = random_split(train_dset, lengths=[55000, 5000])

test_dset = ImageFolder(root="mnist-pngs/test", transform=data_transforms["test"])
```

```
data_transforms = {
    # ...
    "test": transforms.Compose(
        [
            transforms.Resize((32, 32)),
            transforms.CenterCrop((28, 28)),
            transforms.ToTensor(),
            # normalize images to [-1, 1] range
            transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
        ]
    ),
}
```

```
● ● ●

from torch.utils.data import DataLoader

train_loader = DataLoader(
    dataset=train_dset,
    batch_size=batch_size,
    drop_last=True,
    num_workers=4,
    shuffle=True,
)
```



```
from torch.utils.data import DataLoader  
  
train_loader = DataLoader(  
    dataset=train_dset,  
    batch_size=batch_size,  
    drop_last=True,  
    num_workers=4,  
    shuffle=True,  
)
```



```
valid_loader = DataLoader(  
    dataset=valid_dset,  
    batch_size=batch_size,  
    drop_last=False,  
    num_workers=4,  
    shuffle=False,  
)
```



```
from torch.utils.data import DataLoader  
  
train_loader = DataLoader(  
    dataset=train_dset,  
    batch_size=batch_size,  
    drop_last=True,  
    num_workers=4,  
    shuffle=True,  
)
```



```
valid_loader = DataLoader(  
    dataset=valid_dset,  
    batch_size=batch_size,  
    drop_last=False,  
    num_workers=4,  
    shuffle=False,  
)
```



```
test_loader = DataLoader(  
    dataset=test_dset,  
    batch_size=batch_size,  
    drop_last=False,  
    num_workers=4,  
    shuffle=False,  
)
```

Why do I like PyTorch?

It's Pythonic and flexible!

A custom layer

```
import torch

class CoralLayer(torch.nn.Module):
    def __init__(self, size_in, num_classes, preinit_bias=True):
        super().__init__()
        self.size_in, self.size_out = size_in, 1

        self.coral_weights = torch.nn.Linear(self.size_in, 1, bias=False)
        if preinit_bias:
            self.coral_bias = torch.nn.Parameter(
                torch.arange(num_classes - 1, 0, -1).float() / (num_classes-1))
        else:
            self.coral_bias = torch.nn.Parameter(
                torch.zeros(num_classes-1).float())

    def forward(self, x):
        return self.coral_weights(x) + self.coral_bias
```

<https://github.com/Raschka-research-group/coral-pytorch>

A Keras port of the custom layer

```
from typing import Optional
import warnings
import tensorflow as tf
import tensorflow.keras.regularizers

@tf.keras.utils.register_keras_serializable(package="coral_ordinal")
class CoralOrdinal(tf.keras.layers.Layer):

    def __init__(
        self,
        num_classes: int,
        activation: Optional[str] = None,
        kernel_regularizer: Optional[tf.keras.regularizers.Regularizer] = None,
        bias_regularizer: Optional[tf.keras.regularizers.Regularizer] = None,
        **kwargs,
    ):

        if "input_shape" not in kwargs and "input_dim" in kwargs:
            kwargs["input_shape"] = (kwargs.pop("input_dim"),)
        super(CoralOrdinal, self).__init__(**kwargs)
        self.num_classes = num_classes
        self.activation = tf.keras.activations.get(activation)
        self.kernel_regularizer = tf.keras.regularizers.get(kernel_regularizer)
        self.bias_regularizer = tf.keras.regularizers.get(bias_regularizer)

    def build(self, input_shape):
        num_units = 1

        self.kernel = self.add_weight(
            shape=(input_shape[-1], num_units),
            name=self.name + "_latent",
            initializer="glorot_uniform",
            regularizer=self.kernel_regularizer,
            dtype=tf.float32,
            trainable=True,
        )

        self.bias = self.add_weight(
            shape=(self.num_classes - 1,),
            name=self.name + "_bias",
            regularizer=self.bias_regularizer,
            initializer="zeros",
            dtype=tf.float32,
            trainable=True,
        )

    def call(self, inputs):
        kernelized_inputs = tf.matmul(inputs, self.kernel)

        logits = kernelized_inputs + self.bias

        if self.activation is None:
            outputs = logits
        else:
            outputs = self.activation(logits)

        return outputs

    def get_config(self):
        config = super(CoralOrdinal, self).get_config()
        config.update(
            {
                "num_classes": self.num_classes,
                "kernel_regularizer": self.kernel_regularizer,
                "bias_regularizer": self.bias_regularizer,
            }
        )
        return config
```

https://github.com/ck37/coral-ordinal/blob/master/coral_ordinal/layer.py

Live Demo



<https://sebastianraschka.com/books/>
<https://github.com/rasbt/machine-learning-book>



Lightning

D E V ↗ C O N

Get ready to build with Lightning

Join us on June 16th in NYC for the first-ever Lightning Developer Conference

Sign me up!

100% of funds from ticket sales will be used to sponsor travel & expenses for students interested in AI.



<https://www.pytorchlightning.ai/events/devcon2022nyc>